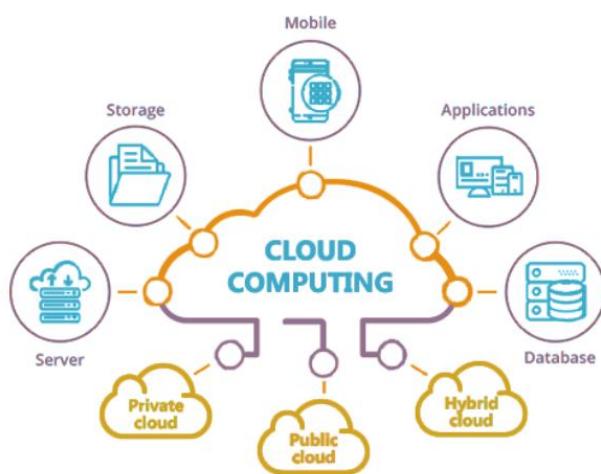


True Engineer

Cloud Computing Notes

Cloud Computing

1. **Cloud:** In the context of computing, the term "cloud" refers to the internet. It's a metaphorical representation of the internet, derived from the cloud symbol often used to depict the internet in network diagrams. So, when we talk about "cloud computing," we're essentially referring to computing services and resources that are delivered over the internet, rather than being hosted on local servers or personal computers.
2. **Computing:** Computing encompasses all activities that involve the use of computers and related technology. This includes processing data, running software applications, storing and retrieving information, and communicating over networks. In the context of cloud computing, "computing" specifically refers to the provision and consumption of computing resources, such as virtual servers, storage, databases, software applications, and other IT services, over the internet.



Cloud computing refers to the delivery of computing services—such as storage, processing power, and software—over the internet. Instead of owning physical hardware or infrastructure, users can access these resources on-demand from a cloud service provider.

Characteristics

There are several key characteristics of cloud computing:

1. **On-Demand Self-Service:** Users can provision computing resources, such as server time and network storage, without human intervention from the service provider.
2. **Broad Network Access:** Services are available over the network and can be accessed through standard mechanisms, promoting use by diverse client platforms (e.g., mobile phones, tablets, laptops).
3. **Resource Pooling:** Cloud resources are pooled to serve multiple consumers, with different physical and virtual resources dynamically assigned and reassigned according to consumer demand.
4. **Rapid Elasticity:** Computing resources can be rapidly and elastically provisioned and released to scale with demand. This ensures that users can access the resources they need when they need them, without worrying about infrastructure limitations.
5. **Measured Service:** Cloud systems automatically control and optimize resource use by leveraging a metering capability at some level of abstraction appropriate to the type of service (e.g., storage, processing, bandwidth, and active user accounts). This allows for transparency and accountability for both the provider and the consumer.

Service Provider's

A service provider in the context of cloud computing refers to a company or organization that offers cloud computing services and resources to individuals, businesses, and other entities. These service providers own and maintain the infrastructure, hardware, and software required to deliver various cloud services over the internet.

There are many types of cloud service providers, ranging from large multinational corporations to smaller, specialized firms. Some of the prominent cloud service providers include:

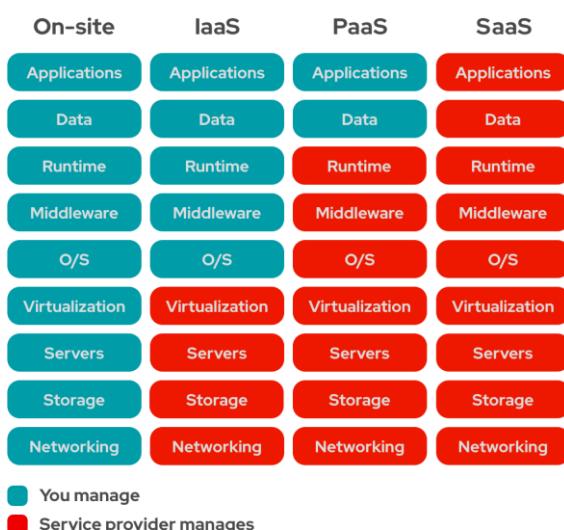
1. **Amazon Web Services (AWS):** AWS is a subsidiary of Amazon.com and is one of the largest and most widely used cloud computing platforms in the world. It offers a wide range of services, including computing power, storage, databases, machine learning, and more.
2. **Microsoft Azure:** Azure is Microsoft's cloud computing platform and services. It provides a variety of cloud services, including virtual computing, storage, databases, AI, and IoT solutions.

TRUE ENGINEER

3. **Google Cloud Platform (GCP)**: GCP is Google's cloud computing platform, offering services such as computing, storage, databases, machine learning, and data analytics.
4. **IBM Cloud**: IBM offers a comprehensive suite of cloud computing services, including infrastructure as a service (IaaS), platform as a service (PaaS), and software as a service (SaaS). It provides solutions for various industries, including healthcare, finance, and manufacturing.
5. **Alibaba Cloud**: Alibaba Cloud is the cloud computing arm of Alibaba Group, offering a wide range of cloud services, including elastic computing, storage, databases, networking, security, and more.

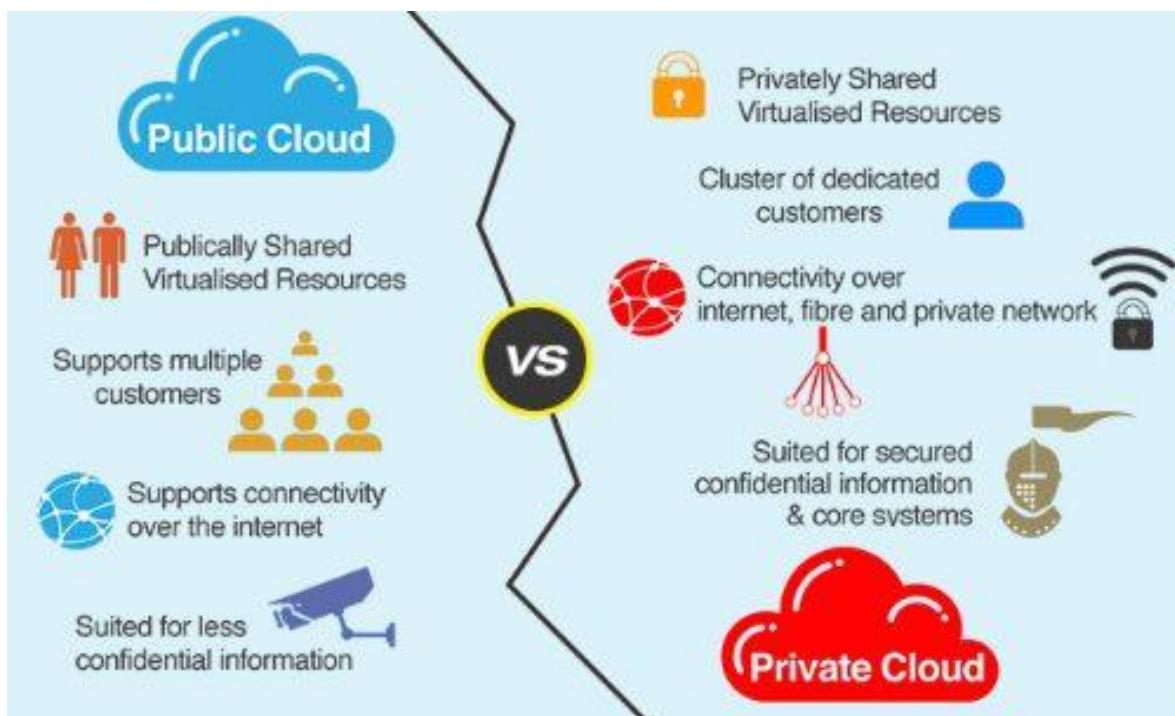
Types of Cloud Services (Service Models)

1. **Infrastructure as a Service (IaaS)**: Provides virtualized computing resources over the internet. Users can rent virtual machines, storage, and networking infrastructure.
2. **Platform as a Service (PaaS)**: Offers a platform allowing customers to develop, run, and manage applications without dealing with the underlying infrastructure. Examples include databases, development tools, and application hosting environments.
3. **Software as a Service (SaaS)**: Delivers software applications over the internet on a subscription basis. Users can access the software through a web browser without needing to install or maintain it locally.



Types Of Cloud (Deployment Models)

- 1. Public Cloud:** Public clouds are owned and operated by third-party cloud service providers, such as Amazon Web Services (AWS), Microsoft Azure, or Google Cloud Platform. These cloud services are made available to the general public or a large industry group and are accessible over the internet. Resources, such as servers and storage, are shared among multiple organizations, resulting in cost savings and scalability benefits.
- 2. Private Cloud:** Private clouds are dedicated cloud environments that are exclusively used by a single organization. They can be hosted on-premises within the organization's data centers or provided by third-party vendors. Private clouds offer greater control, customization, and security compared to public clouds, making them suitable for organizations with specific regulatory requirements or sensitive data.
- 3. Hybrid Cloud:** Hybrid clouds combine elements of both public and private clouds, allowing data and applications to be shared between them. Organizations can use public cloud resources for non-sensitive workloads or to handle spikes in demand, while keeping sensitive workloads and data in a private cloud. Hybrid cloud environments offer flexibility, scalability, and the ability to optimize costs while meeting various business needs.
- 4. Multi-Cloud:** Multi-cloud refers to the use of multiple cloud service providers to host different workloads or services. Organizations may choose a multi-cloud strategy to avoid vendor lock-in, improve resilience and redundancy, optimize performance, or take advantage of best-of-breed services from different providers. Managing a multi-cloud environment can be complex but offers increased flexibility and choice.

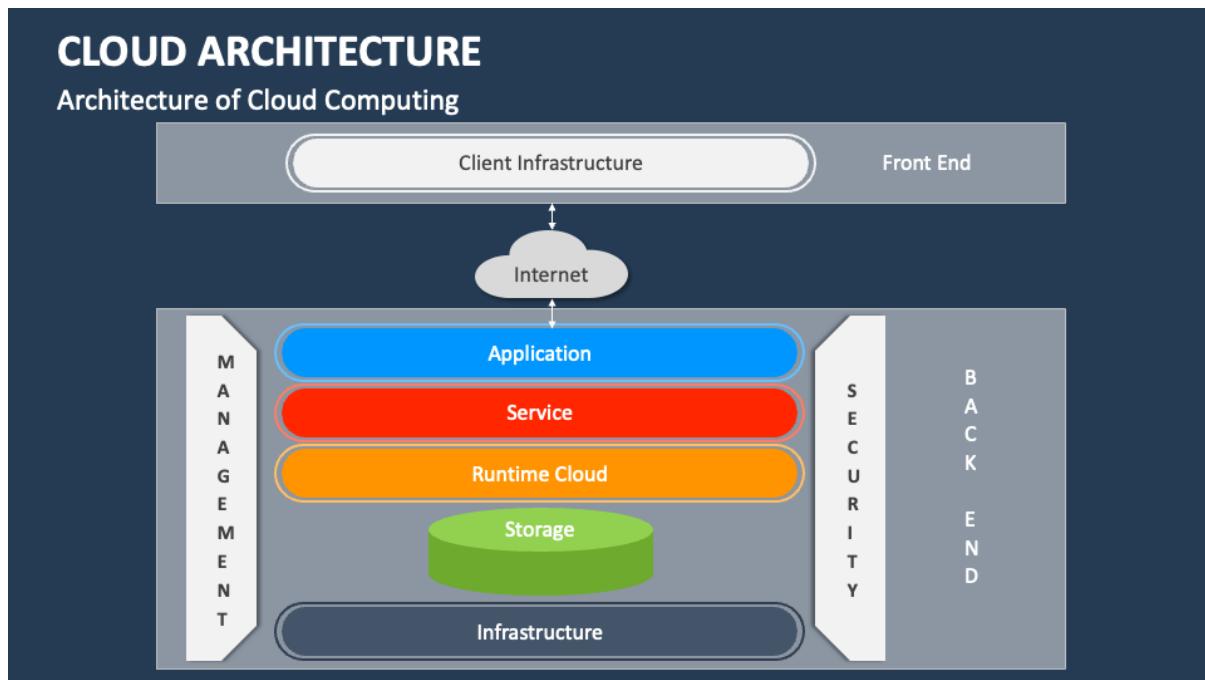


Why Cloud Computing

Cloud computing offers numerous benefits and advantages that make it an attractive choice for individuals, businesses, and organizations of all sizes. Here are some of the key reasons why cloud computing is widely adopted:

- 1. Scalability:** Cloud computing allows users to scale computing resources up or down according to demand. This scalability ensures that organizations can easily handle spikes in traffic or workload without investing in additional hardware or infrastructure.
- 2. Cost-Efficiency:** Cloud computing follows a pay-as-you-go pricing model, where users only pay for the resources they consume. This eliminates the need for upfront investments in hardware and reduces operational costs associated with maintenance, upgrades, and management of on-premises infrastructure.
- 3. Flexibility and Accessibility:** Cloud computing enables users to access computing resources and services from anywhere with an internet connection. This flexibility allows for remote work, collaboration, and access to applications and data on various devices, improving productivity and efficiency.
- 4. Reliability and Redundancy:** Cloud service providers offer high levels of reliability and uptime, with built-in redundancy and failover mechanisms. This ensures that applications and data remain available and accessible even in the event of hardware failures or disasters.
- 5. Security:** Cloud providers invest heavily in security measures to protect data and infrastructure from cyber threats and breaches. They employ encryption, access controls, firewalls, and other security mechanisms to safeguard sensitive information and ensure compliance with industry regulations.
- 6. Automatic Updates and Maintenance:** Cloud providers handle software updates, patches, and maintenance tasks, relieving users of the burden of managing and maintaining infrastructure. This ensures that applications and services remain up-to-date and secure without requiring manual intervention.
- 7. Elasticity:** Cloud computing enables users to dynamically allocate and deallocate computing resources based on changing workload demands. This elasticity allows organizations to efficiently manage resources, optimize performance, and respond quickly to changing business needs.
- 8. Innovation and Agility:** Cloud computing fosters innovation and agility by providing access to a wide range of cutting-edge technologies and services, such as artificial intelligence, machine learning, big data analytics, and Internet of Things (IoT). Organizations can quickly experiment with new ideas and deploy innovative solutions without significant upfront investments or infrastructure constraints.

Cloud Computing Architecture



1. Front End: The front end of a cloud computing system refers to the user interface and client-side components that interact with users or applications. It includes the devices, such as laptops, desktops, tablets, or smartphones, and the software applications or web browsers used to access cloud services. The front end allows users to interact with cloud resources, access data, and execute commands.

2. Back End: The back end of a cloud computing system consists of the cloud infrastructure and resources that provide the computing services. It includes the servers, storage, networking hardware, and other infrastructure components hosted in data centers owned and managed by the cloud service provider. The back end is responsible for processing and storing data, executing applications, and delivering computing services to users over the internet.

3. Cloud-Based Delivery: Cloud-based delivery refers to the delivery model used to provide computing services over the internet. There are three primary service models:

- **Infrastructure as a Service (IaaS):** In IaaS, cloud providers offer virtualized computing resources, such as virtual machines, storage, and networking infrastructure, on a pay-as-you-go basis. Users can deploy and manage virtualized servers and storage resources in the cloud, allowing for flexibility and scalability without the need to invest in physical hardware.

TRUE ENGINEER

- **Platform as a Service (PaaS):** PaaS provides a platform for developing, deploying, and managing applications without the complexity of managing underlying infrastructure. Cloud providers offer development tools, middleware, databases, and other services to support the entire application lifecycle. PaaS enables developers to focus on building and delivering applications without worrying about infrastructure management.
 - **Software as a Service (SaaS):** SaaS delivers software applications over the internet on a subscription basis. Users can access and use applications hosted in the cloud through web browsers or APIs without needing to install or maintain software locally. SaaS providers handle maintenance, updates, and security, allowing users to focus on using the software to meet their business needs.
4. **Network:** The network plays a crucial role in cloud computing architecture, facilitating communication and data transfer between the front end, back end, and cloud-based services. It includes internet connectivity, data center networking infrastructure, and security measures such as firewalls, encryption, and virtual private networks (VPNs) to ensure the confidentiality, integrity, and availability of data transmitted over the network.

Applications of Cloud Computing

1. **Web Hosting and Development:** Cloud computing provides a scalable and cost-effective platform for hosting websites and web applications. Developers can leverage cloud infrastructure and platforms to build, deploy, and manage websites and applications without the need for on-premises hardware.
2. **Data Storage and Backup:** Cloud storage services offer scalable and reliable storage solutions for businesses to store and backup data. Organizations can store large volumes of data in the cloud, access it from anywhere, and ensure data redundancy and disaster recovery.
3. **Big Data Analytics:** Cloud computing provides the computational power and storage capacity required for processing and analyzing large datasets. Organizations can use cloud-based big data analytics platforms to derive insights, make data-driven decisions, and gain a competitive edge.
4. **Artificial Intelligence and Machine Learning:** Cloud computing platforms offer the computational resources and tools needed for training and deploying machine learning models and artificial intelligence algorithms. Businesses can leverage cloud-based AI and ML services for tasks such as image recognition, natural language processing, and predictive analytics.
5. **Software Development and Testing:** Cloud-based development platforms and tools enable developers to collaborate, build, and test software applications more efficiently. Developers can access development environments, version control systems, and testing frameworks in the cloud, reducing development time and costs.

TRUE ENGINEER

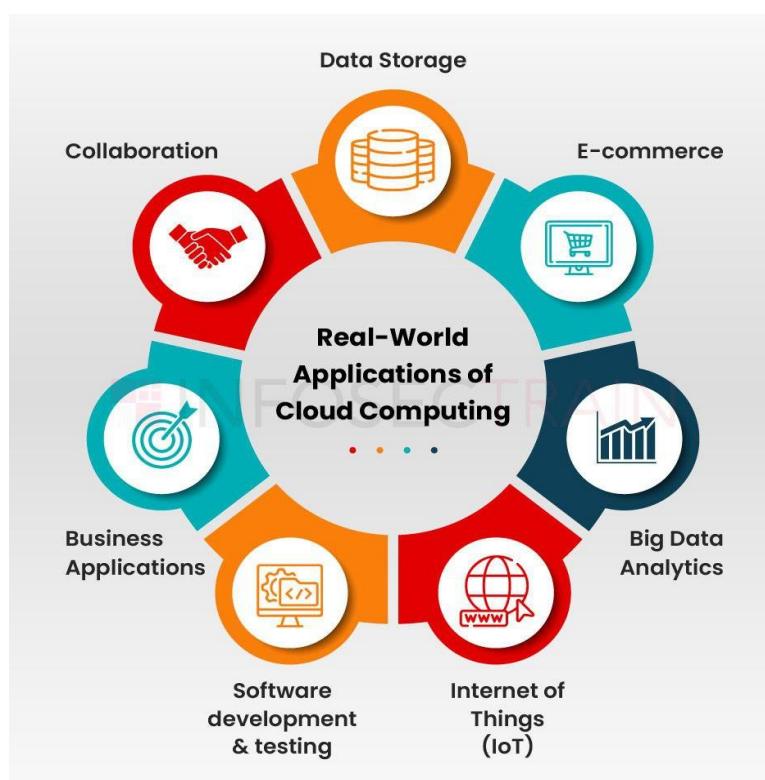
6. Internet of Things (IoT): Cloud computing provides the infrastructure and services required to collect, store, and analyze data generated by IoT devices. Organizations can use cloud-based IoT platforms to process sensor data, monitor devices, and implement IoT applications for various industries, such as healthcare, manufacturing, and smart cities.

7. Enterprise Resource Planning (ERP) and Customer Relationship Management (CRM): Cloud-based ERP and CRM solutions offer businesses centralized platforms for managing core business processes, including finance, HR, sales, and customer service. Cloud-based ERP and CRM systems provide scalability, accessibility, and integration with other business applications.

8. Content Delivery and Media Streaming: Cloud computing enables content providers to deliver digital content, such as videos, music, and games, to users over the internet. Content delivery networks (CDNs) leverage cloud infrastructure to distribute content efficiently, reduce latency, and improve user experience.

9. E-commerce and Online Retail: Cloud computing powers e-commerce platforms and online retail operations, providing scalable and reliable infrastructure for hosting websites, managing inventory, processing transactions, and delivering digital products to customers.

10. Healthcare and Telemedicine: Cloud computing supports healthcare organizations in storing, managing, and analyzing electronic health records (EHRs) and medical imaging data. Telemedicine platforms leverage cloud infrastructure to enable remote consultations, virtual visits, and remote patient monitoring.



Advantages & Dis Advantages of Cloud Computing

Advantages:

1. **Scalability:** Cloud computing allows users to scale computing resources up or down according to demand, enabling organizations to handle spikes in workload without investing in additional hardware.
2. **Cost-Efficiency:** Cloud computing follows a pay-as-you-go pricing model, where users only pay for the resources they consume. This eliminates the need for upfront investments in hardware and reduces operational costs associated with maintenance, upgrades, and management of on-premises infrastructure.
3. **Flexibility and Accessibility:** Cloud computing enables users to access computing resources and services from anywhere with an internet connection. This flexibility allows for remote work, collaboration, and access to applications and data on various devices, improving productivity and efficiency.
4. **Reliability and Redundancy:** Cloud service providers offer high levels of reliability and uptime, with built-in redundancy and failover mechanisms. This ensures that applications and data remain available and accessible even in the event of hardware failures or disasters.
5. **Security:** Cloud providers invest heavily in security measures to protect data and infrastructure from cyber threats and breaches. They employ encryption, access controls, firewalls, and other security mechanisms to safeguard sensitive information and ensure compliance with industry regulations.

Disadvantages:

1. **Dependence on Internet Connectivity:** Cloud computing relies on internet connectivity for accessing resources and services. Organizations may experience downtime or performance issues if they encounter internet connectivity issues or outages.
2. **Data Privacy and Security Concerns:** Storing data in the cloud raises concerns about data privacy, security, and compliance. Organizations must ensure that sensitive data is protected from unauthorized access, breaches, and compliance violations.
3. **Vendor Lock-In:** Adopting cloud-based solutions may lead to vendor lock-in, where organizations become dependent on a single cloud provider for their infrastructure and services. Switching between cloud providers can be challenging and costly due to differences in technologies, APIs, and service offerings.
4. **Limited Control and Customization:** Cloud computing services are managed and controlled by cloud providers, limiting the level of control and customization available to users. Organizations may face limitations in configuring and optimizing their infrastructure and services according to their specific requirements.

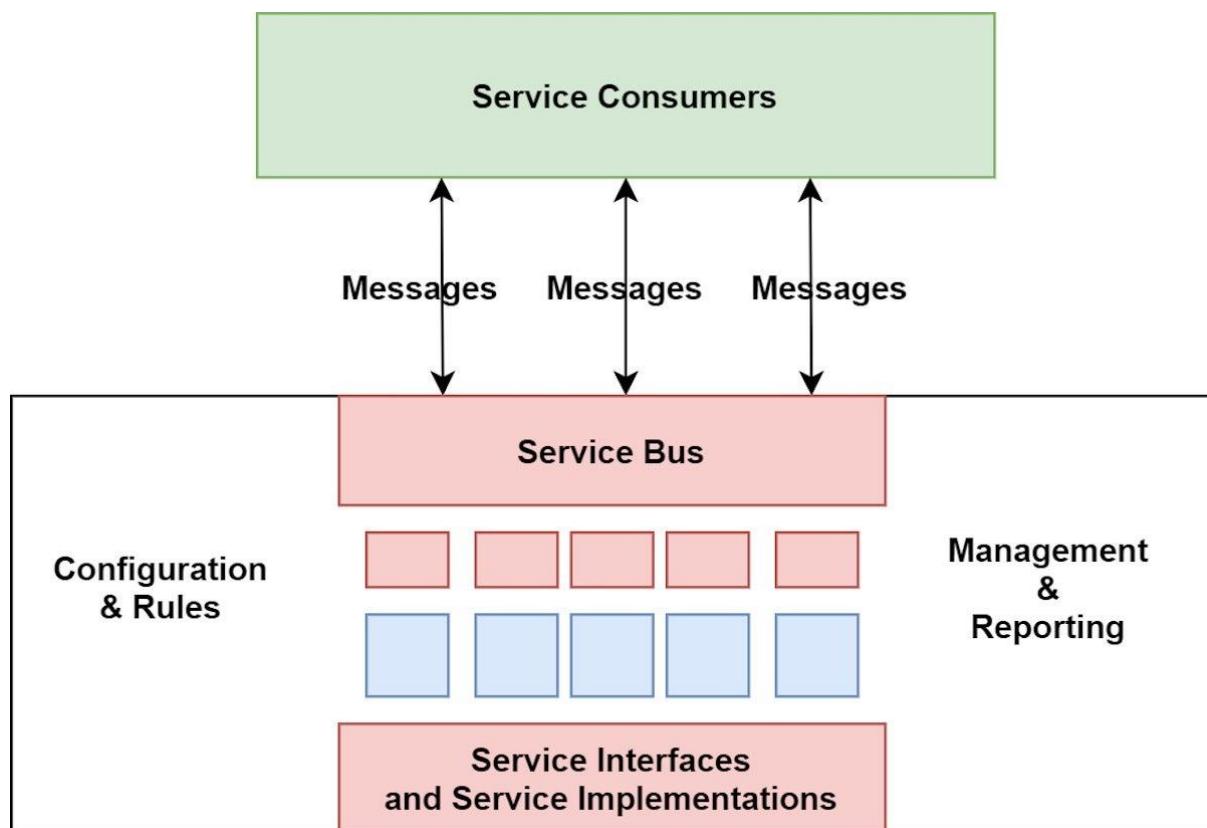
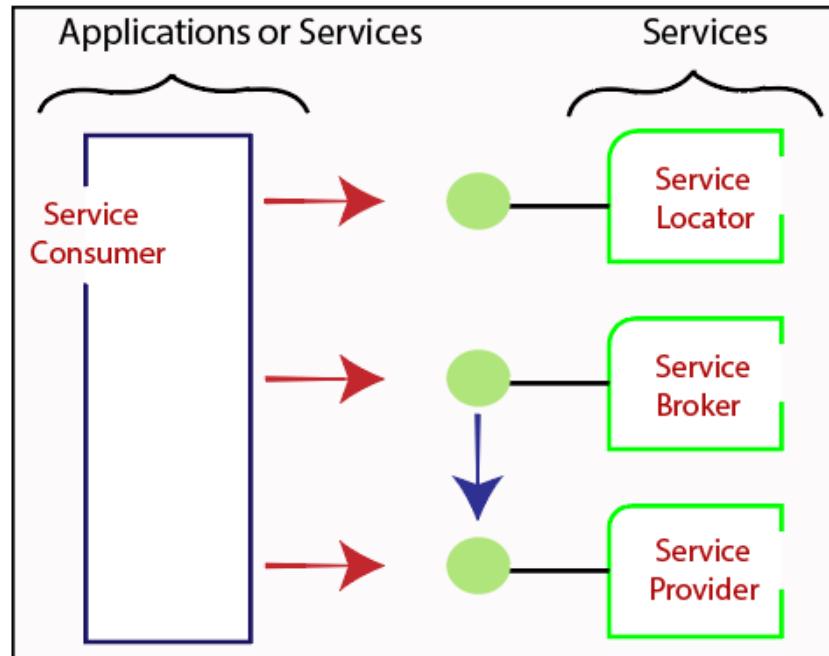
Service-Oriented Architecture (SOA)

Service-oriented architecture (SOA) is a method of software development that uses software components called services to create business applications. Each service provides a business capability, and services can also communicate with each other across platforms and languages. Developers use SOA to reuse services in different systems or combine several independent services to perform complex tasks.

Service-Oriented Architecture (SOA) is an architectural approach to software design that structures an application as a collection of loosely coupled, interoperable services. These services are designed to perform specific business functions and can be accessed and reused across different applications and platforms within an organization or across multiple organizations. SOA promotes modularization, flexibility, and reusability, making it easier to develop, deploy, and maintain complex systems.

Key principles of Service-Oriented Architecture include:

- 1. Loose Coupling:** Services within an SOA are loosely coupled, meaning they are independent and can interact with each other without tight dependencies. This allows services to be developed, deployed, and updated independently, promoting flexibility and agility.
- 2. Interoperability:** SOA promotes interoperability between services, enabling them to communicate and exchange data using standard protocols and interfaces. This allows services to be accessed and reused across different platforms and technologies, facilitating integration and collaboration between systems.
- 3. Service Reusability:** Services within an SOA are designed to be reusable, meaning they can be used in multiple contexts and applications. This reduces redundancy and promotes consistency across systems, improving efficiency and reducing development time and costs.
- 4. Service Abstraction:** SOA encapsulates the underlying implementation details of services, exposing only their interfaces and functionality to consumers. This abstraction allows services to be accessed and used without knowledge of their internal workings, promoting encapsulation and reducing dependencies.
- 5. Service Composition:** SOA enables the composition of services to create larger, more complex applications by orchestrating and coordinating interactions between services. This allows organizations to build flexible and scalable systems by combining existing services to meet specific business needs.
- 6. Service Discoverability:** SOA provides mechanisms for discovering and accessing services within an environment, such as service registries or directories. This allows consumers to easily find and consume services without prior knowledge of their locations or implementations.



What are the benefits of service-oriented architecture?

Service-oriented architecture (SOA) has several benefits over the traditional monolithic architectures in which all processes run as a single unit. Some major benefits of SOA include the following:

Faster time to market

Developers reuse services across different business processes to save time and costs. They can assemble applications much faster with SOA than by writing code and performing integrations from scratch.

Efficient maintenance

It's easier to create, update, and debug small services than large code blocks in monolithic applications. Modifying any service in SOA does not impact the overall functionality of the business process.

Greater adaptability

SOA is more adaptable to advances in technology. You can modernize your applications efficiently and cost effectively. For example, healthcare organizations can use the functionality of older electronic health record systems in newer cloud-based applications.

What are the components in service-oriented architecture?

There are four main components in service-oriented architecture (SOA).

Service

Services are the basic building blocks of SOA. They can be private—available only to internal users of an organization—or public—accessible over the internet to all. Individually, each service has three main features.

Service implementation

The service implementation is the code that builds the logic for performing the specific service function, such as user authentication or bill calculation.

Service contract

The service contract defines the nature of the service and its associated terms and conditions, such as the prerequisites for using the service, service cost, and quality of service provided.

Service interface

In SOA, other services or systems communicate with a service through its service interface. The interface defines how you can invoke the service to perform activities or exchange data. It reduces dependencies between services and the service requester. For example, even users with little or no understanding of the underlying code logic can use a service through its interface.

Service provider

The service provider creates, maintains, and provides one or more services that others can use. Organizations can create their own services or purchase them from third-party service vendors.

Service consumer

The service consumer requests the service provider to run a specific service. It can be an entire system, application, or other service. The service contract specifies the rules that the service provider and consumer must follow when interacting with each other. Service providers and consumers can belong to different departments, organizations, and even industries.

Service registry

A service registry, or service repository, is a network-accessible directory of available services. It stores service description documents from service providers. The description documents contain information about the service and how to communicate with it. Service consumers can easily discover the services they need by using the service registry.

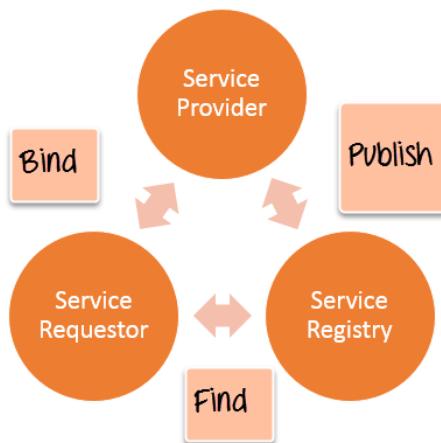
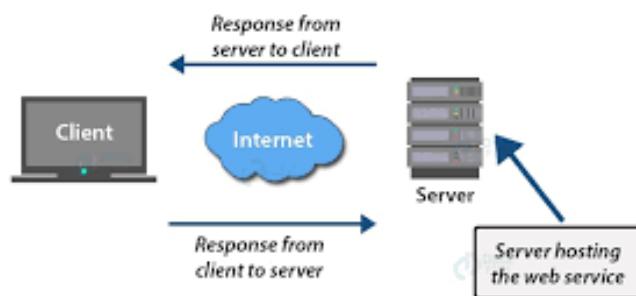
Web Services

In the context of cloud computing, web services refer to software components or applications that provide interoperable communication over the internet using standard protocols and formats.

These services enable different systems and applications to communicate, interact, and exchange data seamlessly, regardless of the underlying platforms, technologies, or programming languages.

Web services in cloud computing typically adhere to the principles of Service-Oriented Architecture (SOA) and are designed to be modular, loosely coupled, and reusable. They expose well-defined interfaces and functionality that can be accessed and invoked by other systems or applications using standard protocols such as SOAP (Simple Object Access Protocol), REST (Representational State Transfer), or JSON (JavaScript Object Notation).

How Web Servers Work?



Examples

Examples of web services in cloud computing include:

- **SOAP-based Web Services:** SOAP (Simple Object Access Protocol) is a protocol for exchanging structured information in the implementation of web services. SOAP-based web services use XML (Extensible Markup Language) for message exchange and typically follow a contract-based approach using WSDL (Web Services Description Language) for defining service interfaces.
- **RESTful APIs:** REST (Representational State Transfer) is an architectural style for designing networked applications, often used for building lightweight, scalable web services. RESTful APIs use standard HTTP methods (GET, POST, PUT, DELETE) for data exchange and typically represent resources using URIs (Uniform Resource Identifiers) and JSON (JavaScript Object Notation) for data formatting.

Characteristics

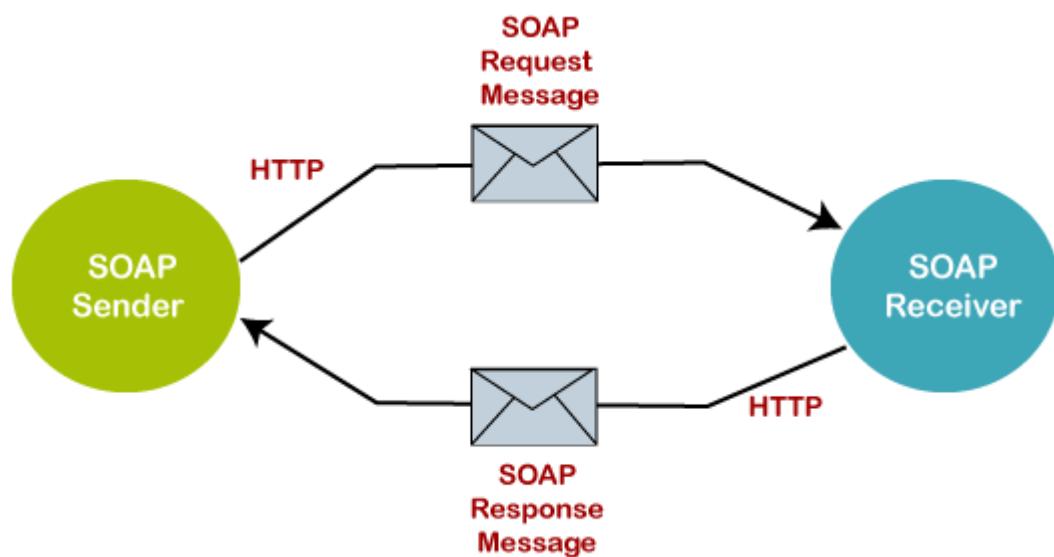
1. **Interoperability:** Web services facilitate interoperable communication between different systems and applications, enabling them to exchange data and invoke functionality regardless of their underlying technologies or platforms.
2. **Loose Coupling:** Web services are loosely coupled, meaning they are independent and can interact with each other without tight dependencies. This promotes flexibility, scalability, and reusability in building and integrating distributed systems.
3. **Standardized Interfaces:** Web services expose standardized interfaces and protocols for accessing and invoking their functionality, making it easier for developers to integrate and interact with them across different platforms and programming languages.
4. **Platform-Independence:** Web services are platform-independent, meaning they can be accessed and invoked from any platform or device with internet connectivity. This enables organizations to build and deploy distributed systems that are accessible from anywhere, promoting flexibility and accessibility.
5. **Scalability and Performance:** Web services can be deployed and scaled to handle varying levels of workload and demand, ensuring optimal performance and responsiveness even under high traffic conditions.
6. **Security:** Web services in cloud computing adhere to security best practices and standards to ensure the confidentiality, integrity, and availability of data and services. This includes authentication, authorization, encryption, and other security mechanisms to protect against unauthorized access and data breaches.

SOAP (Simple Object Access Protocol):

SOAP is a protocol for exchanging structured information in the implementation of web services. It provides a standardized format for sending and receiving messages between distributed applications over the internet or other network protocols.

Key Characteristics of SOAP:

- 1. Message Format:** SOAP messages are typically formatted using XML (Extensible Markup Language), making them human-readable and platform-independent. A SOAP message consists of an envelope containing a header and a body, which may include the actual payload or data being exchanged between the sender and receiver.
- 2. Protocol Independence:** SOAP messages can be transported over various network protocols, including HTTP, SMTP, and TCP/IP. This protocol independence allows SOAP-based web services to communicate across different systems and platforms.
- 3. Contract-Based Communication:** SOAP-based web services often follow a contract-based approach using WSDL (Web Services Description Language) for defining service interfaces. WSDL specifies the operations, messages, and data types supported by the web service, enabling clients to discover and invoke service methods dynamically.
- 4. Support for Security:** SOAP supports various security mechanisms, including encryption, digital signatures, and authentication, to ensure the confidentiality, integrity, and authenticity of messages exchanged between clients and services.
- 5. Complexity and Overhead:** SOAP messages can be more complex and have higher overhead compared to other web service protocols, such as REST. This complexity arises from the XML-based message format and additional protocol features, such as SOAP headers and fault handling.



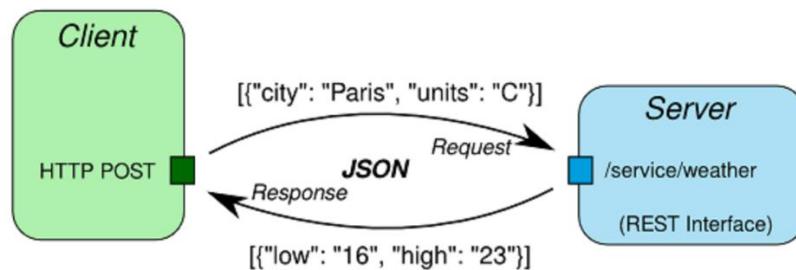
REST (Representational State Transfer):

REST is an architectural style for designing networked applications, particularly web services, that emphasizes simplicity, scalability, and flexibility. RESTful web services use standard HTTP methods (GET, POST, PUT, DELETE) for data exchange and represent resources using URIs (Uniform Resource Identifiers) and various media types, such as JSON (JavaScript Object Notation) or XML.

Key Characteristics of REST:

- 1. Resource-Based Architecture:** RESTful web services model resources as entities that can be identified by unique URIs. Resources represent real-world entities, such as users, products, or orders, and are manipulated using standard HTTP methods (GET, POST, PUT, DELETE) to perform CRUD (Create, Read, Update, Delete) operations.
- 2. Statelessness:** RESTful web services are stateless, meaning each request from a client to the server contains all the information necessary to process the request. This allows services to scale easily and be deployed across distributed environments without relying on server-side state or session management.
- 3. Uniform Interface:** RESTful web services use a uniform interface for communication, consisting of standard HTTP methods (GET, POST, PUT, DELETE) and resource representations (e.g., JSON, XML). This simplifies client-server interactions and promotes interoperability between systems and platforms.
- 4. Cacheability:** RESTful web services support caching of resources to improve performance and scalability. Responses from the server can be cached by clients or intermediate proxies based on cache-control headers, reducing the need for repeated requests to the server.
- 5. Lightweight and Scalable:** RESTful web services are lightweight and have lower overhead compared to SOAP-based services, making them well-suited for distributed and cloud-based architectures. They can be easily consumed by a wide range of clients, including web browsers, mobile devices, and IoT (Internet of Things) devices.

RESTful Web Service in Java



Web Services Description Language

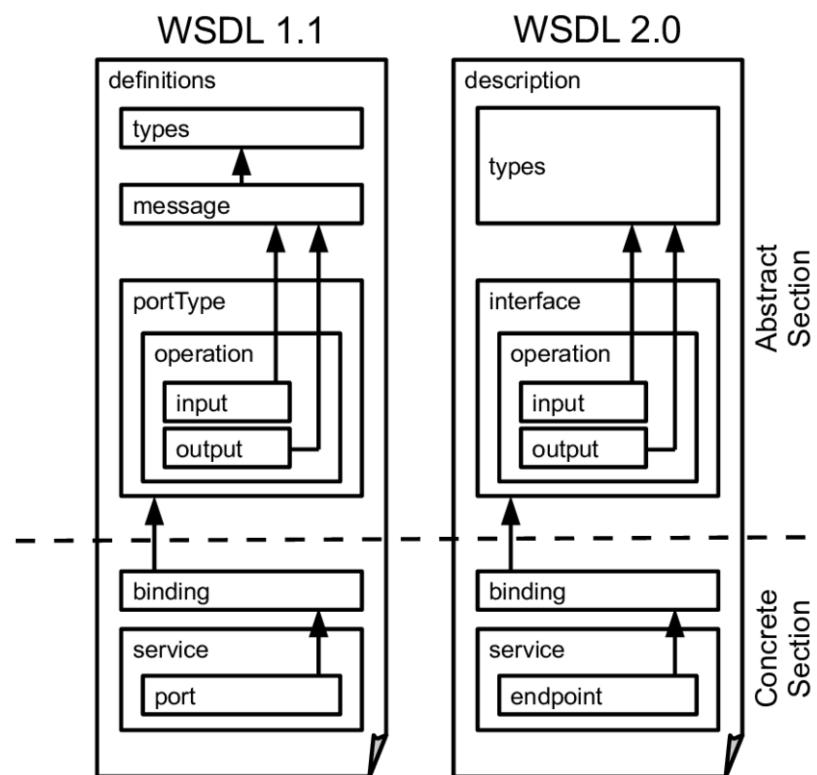
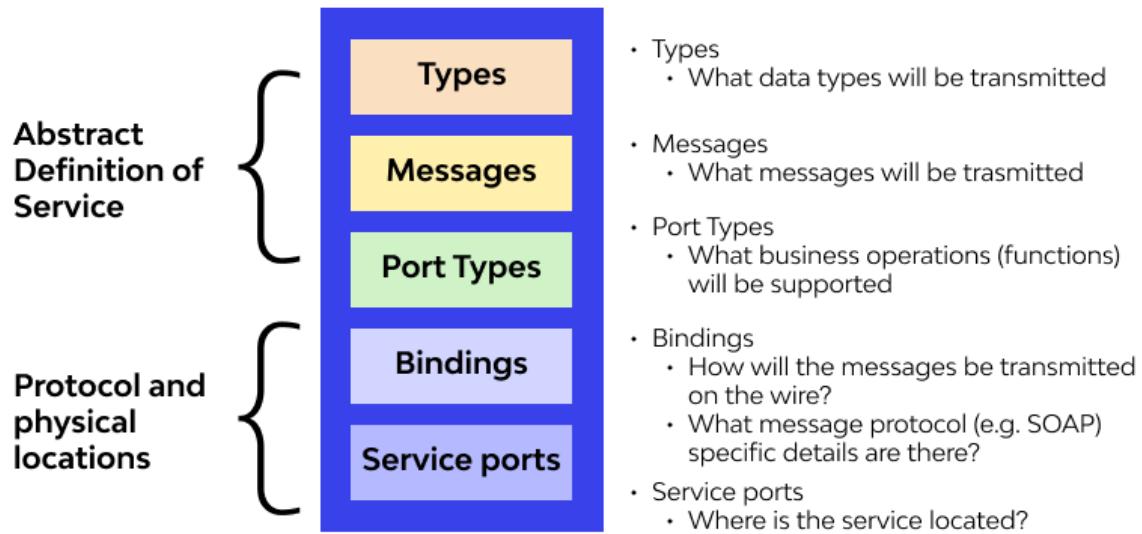
Web Services Description Language (WSDL) is an XML-based language used to describe the functionality offered by a web service. It provides a standardized way for service providers to define the interfaces, operations, and data types supported by their web services, allowing clients to discover, understand, and interact with the services dynamically.

Key Components of WSDL:

1. **Types:** The `<types>` element in WSDL defines the data types used by the web service, such as primitive types (e.g., integer, string) and complex types (e.g., structures, arrays). Data types are typically defined using XML Schema (XSD), allowing for interoperability between different systems and platforms.
2. **Message:** The `<message>` element in WSDL describes the abstract message formats exchanged between clients and the web service. Messages are composed of one or more parts, each representing a parameter or piece of data associated with a particular operation.
3. **Port Type:** The `<portType>` element in WSDL defines the operations supported by the web service and the messages associated with each operation. Each operation specifies the input and output messages it expects or produces, allowing clients to invoke the operations and exchange data with the service.
4. **Binding:** The `<binding>` element in WSDL specifies the communication protocols and message formats used by the web service. It defines how the abstract operations defined in the `<portType>` element are mapped to concrete network protocols, such as SOAP over HTTP or RESTful HTTP.
5. **Service:** The `<service>` element in WSDL identifies the endpoints and locations where the web service is available. It contains one or more `<port>` elements, each representing a specific endpoint or network address where clients can access the service.

WSDL Elements

A WSDL document describes a web service using these major elements:



Hypervisor

A hypervisor is a software that you can use to run multiple virtual machines on a single physical machine. Every virtual machine has its own operating system and applications. The hypervisor allocates the underlying physical computing resources such as CPU and memory to individual virtual machines as required. Thus, it supports the optimal use of physical IT infrastructure.

A hypervisor, also known as a virtual machine monitor (VMM), is a software or firmware layer that enables multiple virtual machines (VMs) to run on a single physical machine, known as a host. The hypervisor abstracts the underlying hardware resources of the host, such as CPU, memory, storage, and networking, and allocates them to virtual machines, allowing each VM to operate as if it were running on its own dedicated physical hardware.

Key characteristics of a hypervisor include:

- 1. Virtualization:** The hypervisor creates and manages virtualized environments, known as virtual machines (VMs), which share the physical resources of the host system. Each VM runs its own operating system and applications, isolated from other VMs running on the same host.
- 2. Resource Management:** The hypervisor allocates and manages hardware resources, such as CPU, memory, storage, and networking, among multiple VMs running on the host. It ensures fair and efficient resource utilization, prioritizing critical workloads and preventing resource contention.
- 3. Isolation:** The hypervisor provides strong isolation between virtual machines, ensuring that each VM operates independently and securely without interference from other VMs running on the same host. This isolation prevents one VM from affecting the performance or stability of other VMs.
- 4. Hardware Abstraction:** The hypervisor abstracts the underlying hardware of the host system, presenting a virtualized view of the hardware resources to each VM. This abstraction enables VMs to run different operating systems and applications without being aware of the underlying physical hardware.
- 5. Live Migration:** Some hypervisors support live migration, allowing VMs to be moved from one physical host to another without disruption to running applications or services. Live migration enables workload balancing, resource optimization, and fault tolerance in virtualized environments.
- 6. Management Interfaces:** Hypervisors typically provide management interfaces, such as command-line tools, graphical user interfaces (GUIs), and application programming interfaces (APIs), for administering and monitoring virtualized environments. These interfaces allow administrators to create, configure, monitor, and manage VMs and their resources.

Why is a hypervisor important?

Hypervisors are the underlying technology behind virtualization or the decoupling of hardware from software. IT administrators can create multiple virtual machines on a single host machine. Each virtual machine has its own operating system and hardware resources such as a CPU, a graphics accelerator, and storage. You can install software applications on a virtual machine, just like you do on a physical computer.

The fundamentals of virtual machines and other virtualization technologies have enabled cloud computing services in enterprise applications. They allow you to scale computing services efficiently on limited hardware infrastructure. For example, different business departments can run different workloads separately by using multiple virtual machines on a single server.

What are the benefits of a hypervisor?

Organizations use virtualization software like hypervisors because the software helps them to use resources efficiently and reduce hardware investment. Virtualization brings several other benefits such as those given below.

Hardware independence

A hypervisor abstracts the host's hardware from the operating software environment. IT administrators can configure, deploy, and manage software applications without being constrained to a specific hardware setup. For example, you can run macOS on a virtual machine instead of iMac computers.

Efficiency

Hypervisors make setting up a server operating system more efficient. Manually installing the operating system and related software components is a time-consuming process. Instead, you can configure the hypervisor to immediately create your virtual environment.

Scalability

Organizations use hypervisors to maximize resource usage on physical computers. Instead of using separate machines for different workloads, hypervisors create multiple virtual computers to run several workloads on a single machine. This translates to faster scalability and reduced hardware expenditure for organizations.

Portability

IT teams can allocate memory, networking, processing, and storage resources across multiple servers as needed. They have the ability to shift workloads between machines or platforms easily. When an application requires more processing power, the hypervisor provides seamless access to additional physical resources.

What are the use cases for hypervisors?

Virtualization software that are powered by hypervisors have several use cases. We give some examples below.

Desktop virtualization

Employees use desktop virtualization software to emulate a version of their workstation computing environment on the server. This allows them to access their work applications and files remotely.

Resource optimization

Companies use hypervisors to consolidate multiple computers that perform different functions into a single server. For example, if production, marketing, and customer support teams run their workloads on individual physical servers, it can result in idle resources. With a hypervisor, you can host the virtual machines for the respective business units on a single server, even if they require different operating systems and software components.

Failure recovery

The hypervisor captures snapshots of the virtual machine's previous state in a virtual machine image—a file that contains the installation instructions, configurations, and other details of the virtual machine. System administrators can use the image file to restore the virtual machine in case of failure. There is also capability to create backup copies or move the virtual machine to a different host.

Legacy system continuity

Some organizations have invested significantly in software that has outlasted the underlying server. Hypervisors provide an option to continue running the software by virtualizing the hardware environment required. This allows organizations to support their cloud transformation efforts with minimum disruption to existing business workflows.

What are the types of hypervisors?

There are two types of hypervisors, each differing in architecture and performance.

Type 1 hypervisor

The type 1 hypervisor sits on top of the metal server and has direct access to the hardware resources. Because of this, the type 1 hypervisor is also known as a bare-metal hypervisor. The host machine does not have an operating system installed in a bare-metal hypervisor setup. Instead, the hypervisor software acts as a lightweight operating system.

Pros and cons

Due to its architecture, the type 1 hypervisor is very efficient. It can directly manage and allocate resources for multiple virtual machines without going through the host operating system. These types of hypervisors are also more secure, as the absence of a host operating system reduces the risks of instability.

Type 2 hypervisor

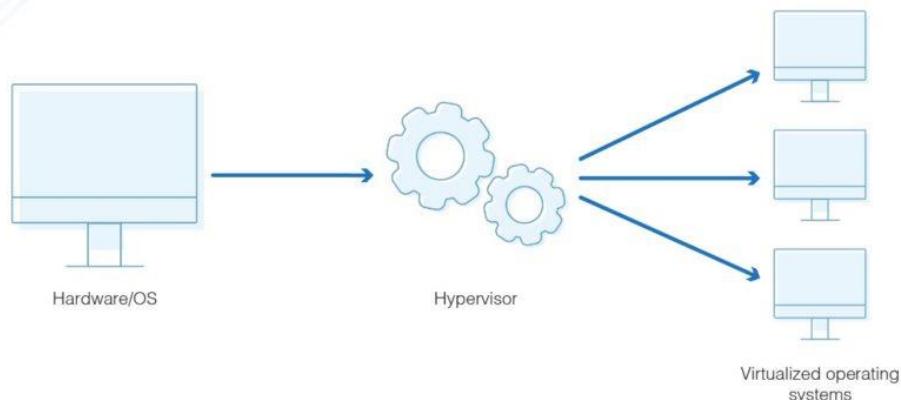
The type 2 hypervisor is a hypervisor program installed on a host operating system. It is also known as a hosted or embedded hypervisor. Like other software applications, hosted hypervisors do not have complete control of the computer resources. Instead, the system administrator allocates the resources for the hosted hypervisor, which it distributes to the virtual machines.

Pros and cons

The presence of the host operating system introduces latency to the virtualized environment. When the virtual machine requests computing resources, the hypervisor cannot directly access the underlying hardware but relays the request to the host operating system. Also, the hypervisor and its hosted virtual machines are dependent on the stability of the host operating system.

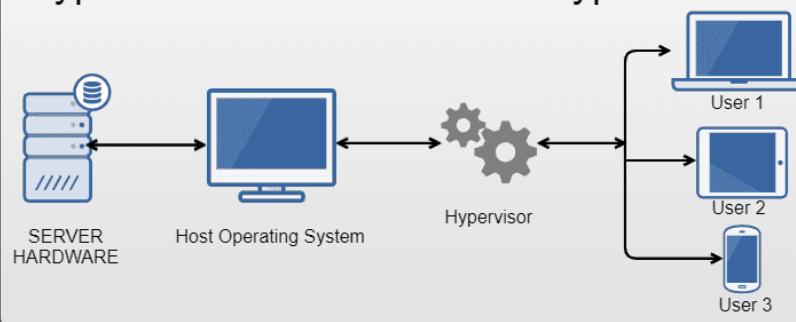
Fig Of Hypervisor

What Is a Hypervisor?

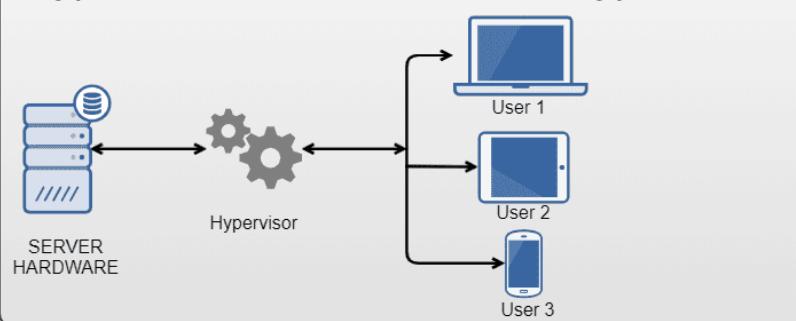


Type of Hypervisor

Type I / Bare-metal / Native Hypervisor



Type II / Embedded / Hosted Hypervisor



Unit 02

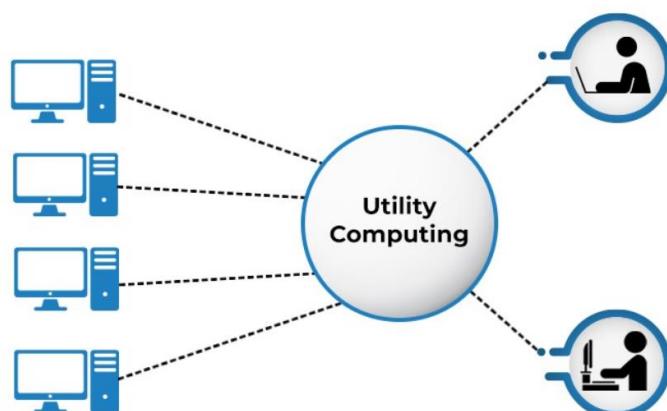
Utility Computing

Utility computing is a service provisioning model that offers computing resources such as hardware, software, and network bandwidth to clients as and when they require them on an on-demand basis. The service provider charges only as per the consumption of the services, rather than a fixed charge or a flat rate.

Utility computing is a model where computing resources, such as processing power, storage, and applications, are provided to users on demand, much like a traditional utility such as electricity or water. Users typically pay for these resources on a metered basis, meaning they are charged for the actual amount of resources they consume, rather than a flat fee.

This model offers several advantages, including scalability, flexibility, and cost-efficiency. Users can easily scale their computing resources up or down based on their current needs, without having to invest in and manage their own infrastructure. Additionally, users can access these resources from anywhere with an internet connection, making it particularly attractive for businesses with dynamic or unpredictable computing needs.

Utility Computing



Utility computing is a subset of [cloud computing](#), allowing users to scale up and down based on their needs. Clients, users, or businesses acquire amenities such as data storage space, computing capabilities, applications services, virtual servers, or even hardware rentals such as CPUs, monitors, and input devices.

The utility computing model is based on conventional utilities and originates from the process of making IT resources as easily available as traditional public utilities such as electricity, gas, water, and telephone services. For example, a consumer pays his electricity bill as per the number of units consumed, nothing more and nothing less. Similarly, utility computing works on the same concept, which is a pay-per-use model.

The service provider owns and manages the computing solutions and infrastructure, and the client subscribes to the same and is charged in a metered manner without any upfront cost. The concept of utility computing is simple—it provides processing power when you need it, where you need it, and at the cost of how much you use it.

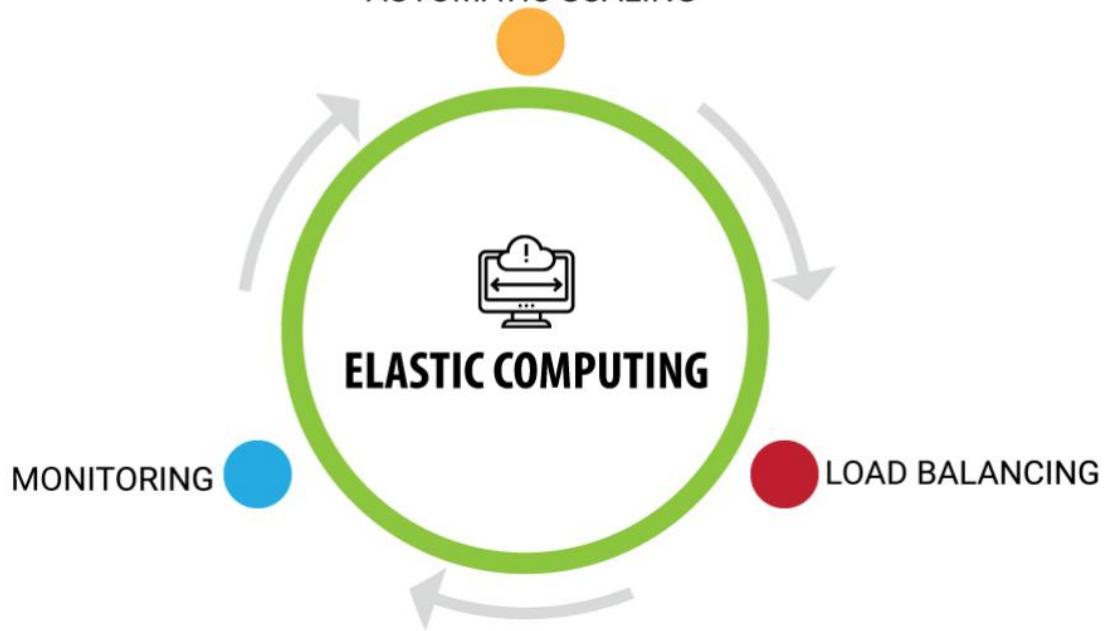
Elastic Computing or Elastic Cloud Computing (EC2)

Elastic computing is the ability of a cloud service provider to swiftly scale the usage of resources such as storage, infrastructure, computer processing, CPU memory, RAM, input/output bandwidth, etc., up and down to adapt to changing resource demands and dynamically meet workload requirements.

Elastic computing is a part of cloud computing that entails dynamically managing the cloud server.

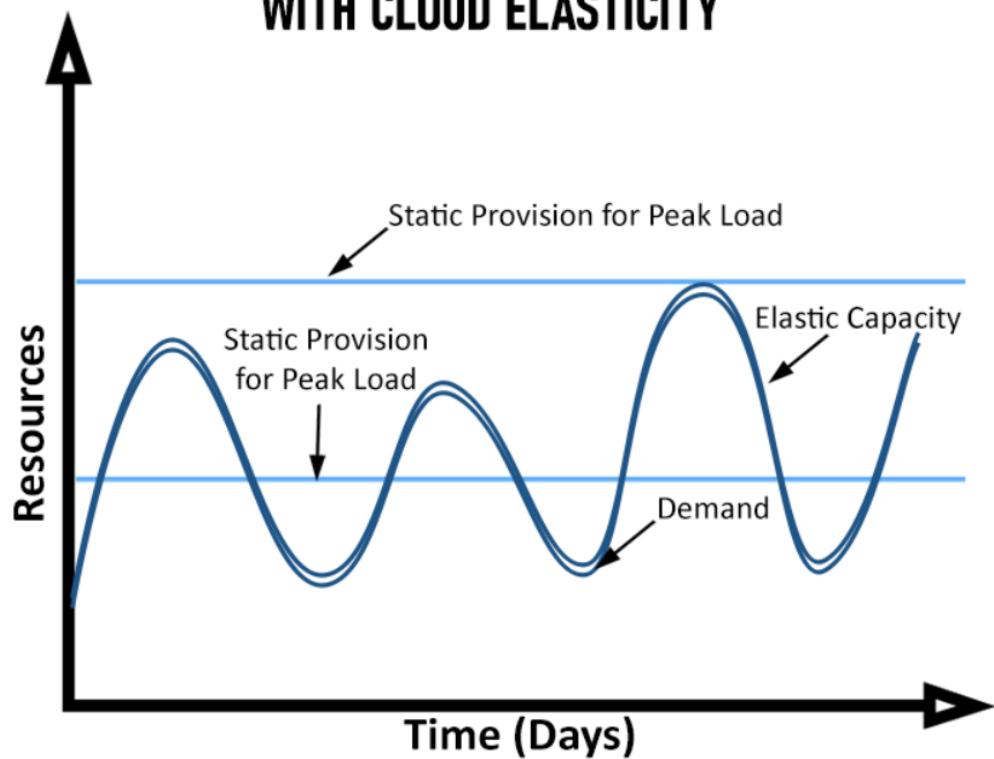
TRUE ENGINEER

AUTOMATIC SCALING



How Elastic Computing Works

COMPARISON OF STATIC CAPACITY WITH CLOUD ELASTICITY



TRUE ENGINEER CLOUD COMPUTING NOTES

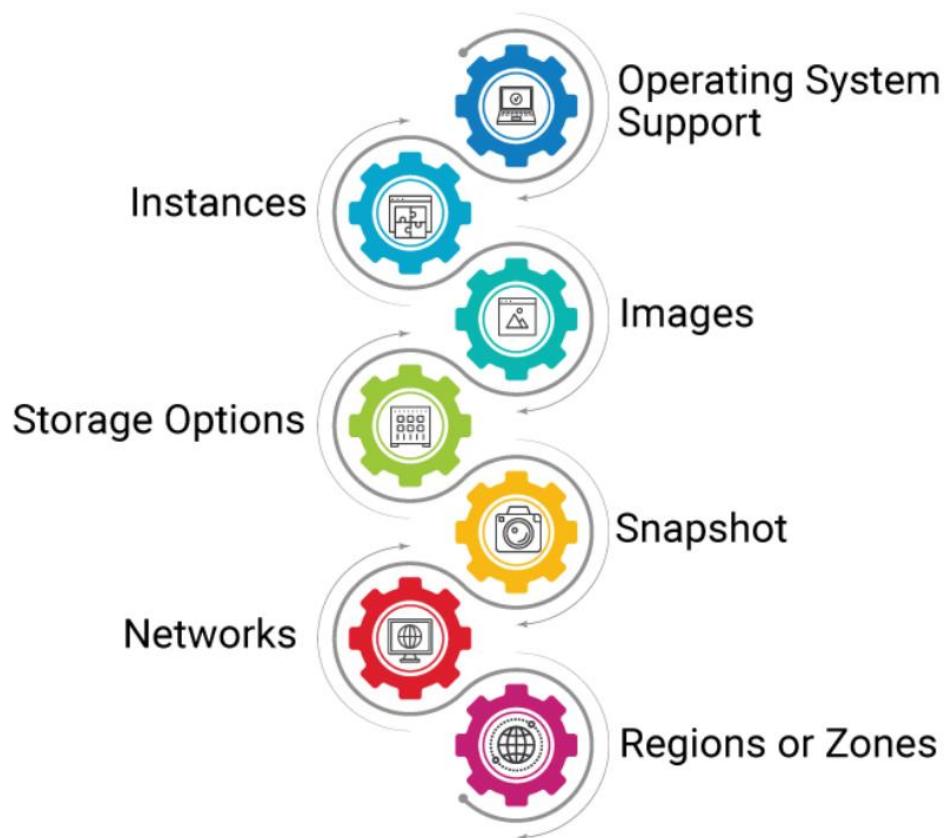
TRUE ENGINEER

Elastic computing is a concept closely related to utility computing and cloud computing. It refers to the ability to dynamically adjust the amount of computing resources allocated to a workload based on changing demands.

In an elastic computing environment, resources such as processing power, memory, and storage can be automatically scaled up or down in response to fluctuations in workload or user demand. This elasticity allows organizations to efficiently use resources, ensuring that they have enough capacity to handle peak loads without over-provisioning and wasting resources during periods of lower demand.

Elasticity is a key feature of cloud computing platforms, where resources can be provisioned and de-provisioned automatically through tools like auto-scaling policies. This flexibility enables businesses to optimize costs, improve performance, and maintain high availability for their applications and services.

KEY COMPONENTS OF ELASTIC COMPUTING



Benefits of elastic computing

Elastic computing is witnessing innovations at an increasingly rapid pace and is driving digital transformation across the IT sector. Let's look at the benefits of elastic computing and its effects on organizations today.

1. Pay-per-use feature

The compelling pay-for-what-you-use feature of elastic cloud computing is one of the main reasons for its growing popularity and massive adoption rate. What separates elastic computing from on-premise computing is that you pay only for the resources you're using, not a penny more. This is unlike on-premise computing, where you keep paying the same amount no matter how many resources you use. This helps organizations save a significant amount that was previously being unnecessarily spent on idle resources.

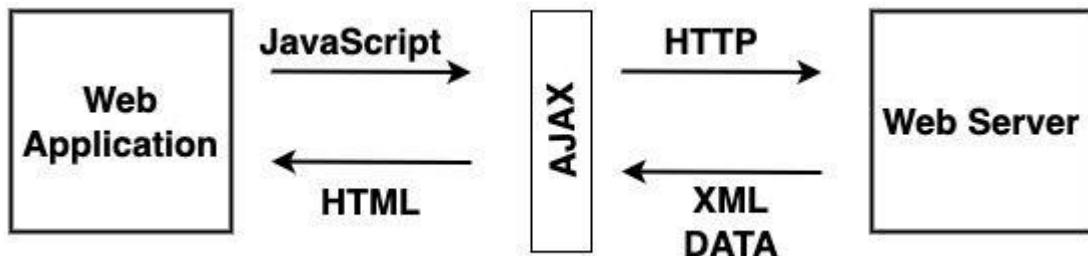
2. Cost-effectiveness & flexibility

Elastic computing enables organizations to avoid huge costs and expenses related to adding extra resources to their data centers. With elastic computing, you no longer have to pay additional costs for unutilized capacity. Formerly, businesses had to purchase more resources even if they didn't require them every day, just to be prepared for an unexpected spike in demand. However, the pay-as-you-go model has solved this problem, wherein you're only charged for the capacity that you actually use.

3. On-demand computing

Gone are the days when surges in website traffic or bandwidth spikes were solved by expanding the architecture with additional servers. The problem with this effort was that it often required thorough capacity planning that needed to be done months in advance, along with enduring enormous up-front costs spent on setting up hardware. Businesses that leverage the advantages of elastic computing no longer have to worry about capacity since it's already in place and completely ready to scale up or down as and when required.

AJAX



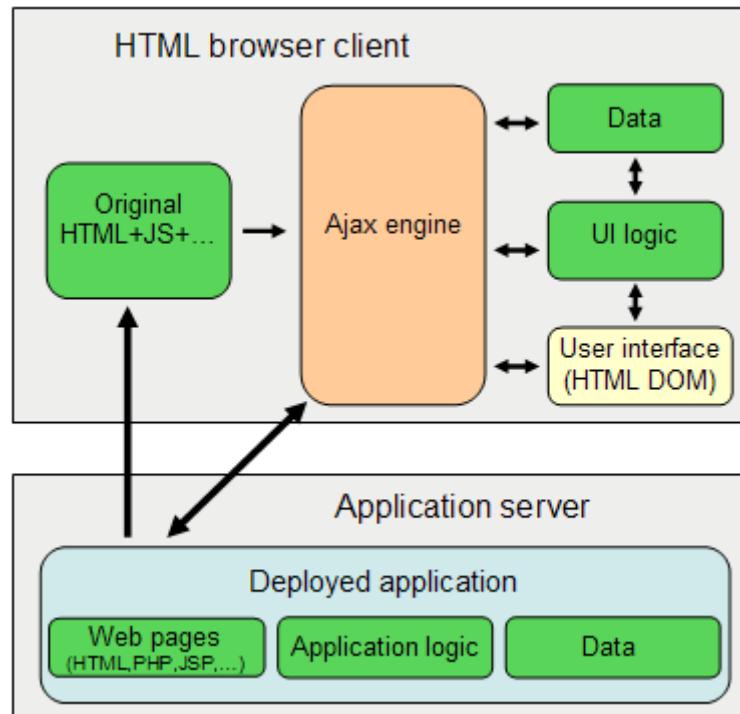
Ajax, which stands for Asynchronous JavaScript and XML, is a set of web development techniques used to create interactive web applications. It allows web pages to send and receive data from a server asynchronously, without requiring the page to be reloaded. This enables more dynamic and responsive user experiences, similar to desktop applications.

Ajax typically involves the use of several technologies working together:

1. JavaScript: Ajax relies heavily on JavaScript to make asynchronous requests to the server and handle the response.
2. XML or JSON: While XML was initially used for data interchange in Ajax, JSON (JavaScript Object Notation) has become more popular due to its lightweight and easier integration with JavaScript.
3. XMLHttpRequest (XHR) object: This is a built-in browser object used to make HTTP requests from JavaScript. It forms the backbone of Ajax functionality.

Ajax is commonly used for various purposes, such as:

- Loading new content dynamically without refreshing the entire page.
- Submitting form data to the server and updating parts of the page based on the response.
- Implementing autocomplete or live search functionality.
- Fetching data from a server to update charts, tables, or other visual components on a webpage.



Simple Definition

Asynchronous JavaScript and XML (AJAX) is a combination of web application development technologies that make web applications more responsive to user interaction. Whenever your users interact with a web application, such as when they click buttons or checkmark boxes, the browser exchanges data with the remote server. Data exchange can cause pages to reload and interrupt the user experience. With AJAX, web applications can send and receive data in the background so that only small portions of the page refresh as required.

Use Cases

Autocomplete

Search engines provide autocomplete options in real time when users search for a specific keyword in the search bar. AJAX allows the webpage to relay each character input to the web server and return a list of relevant recommendations on the existing page.

Form verification

AJAX allows web applications to validate specific information in forms before users submit them. For example, when a new user creates an account, the webpage can automatically verify if a username is available before the user moves to the next section.

Chat functionality

Text messengers and chatbots use AJAX to display real-time conversations on browsers. AJAX sends the text written by a user to the server and publishes it simultaneously in other users' chat interfaces.

Social media

Social media platforms use AJAX to update users' feeds with the latest content without loading a new page on the browser. For example, Twitter refreshes your feed immediately whenever someone you follow tweets an update.

Voting and rating systems

Some forums and social bookmarking sites use AJAX to display the rating or votes of specific posts in real time. For example, you can upvote or downvote a post on Reddit without refreshing the entire page.

How does AJAX work?

AJAX uses JavaScript and XML to enable asynchronous calls when browsers and servers exchange data. Next, we explain how browsers traditionally exchange data and compare it to data exchange with AJAX.

Data exchange without AJAX

In a conventional model, the browser sends an HTTP request to the server side when the user performs an action. The web server receives and processes the request and sends the updated data to the browser. Then, the browser refreshes the webpage with the new data.

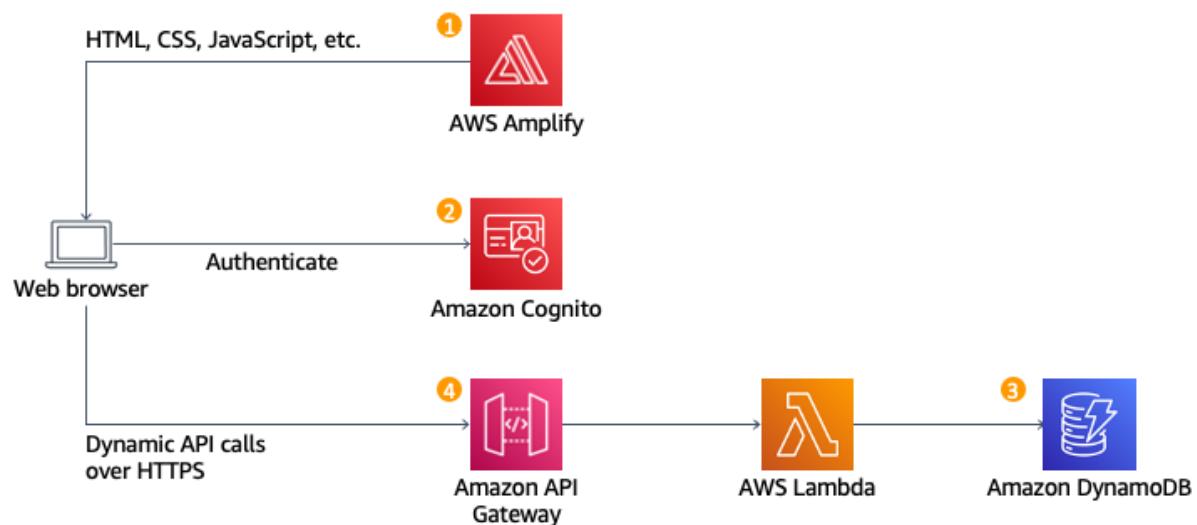
In this approach, the browser reloads the entire page even if the requested data consists of minor changes. Moreover, the browser might send frequent requests, which load the web server software.

Data exchange with AJAX

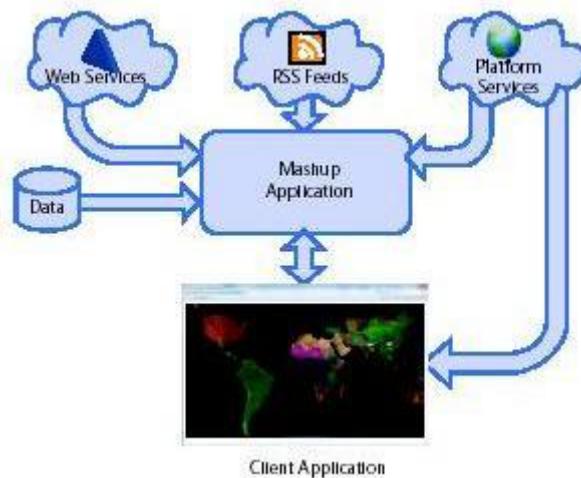
Instead of updating the whole page, AJAX uses a JavaScript function to create an XMLHttpRequest object on the browser. Then, it compiles the page information in XML format, which the XMLHttpRequest object sends to the web server. The web server processes the request and responds with the requested data. Lastly, the browser updates the current screen with the latest data without refreshing the page.

Why is AJAX more efficient?

Despite similarities in data exchange and information flow, AJAX is more efficient than conventional web requests. With AJAX, the browser only updates specific web content based on the requested data. It doesn't make unnecessary refreshes on other content on the page. This makes AJAX applications faster and more responsive than conventional web applications.



Mashups



In the realm of cloud computing, "mashups" refer to the integration of multiple web services or APIs to create a new application or service. This concept stems from the idea of combining various data sources or functionalities to provide users with a more comprehensive or customized solution.

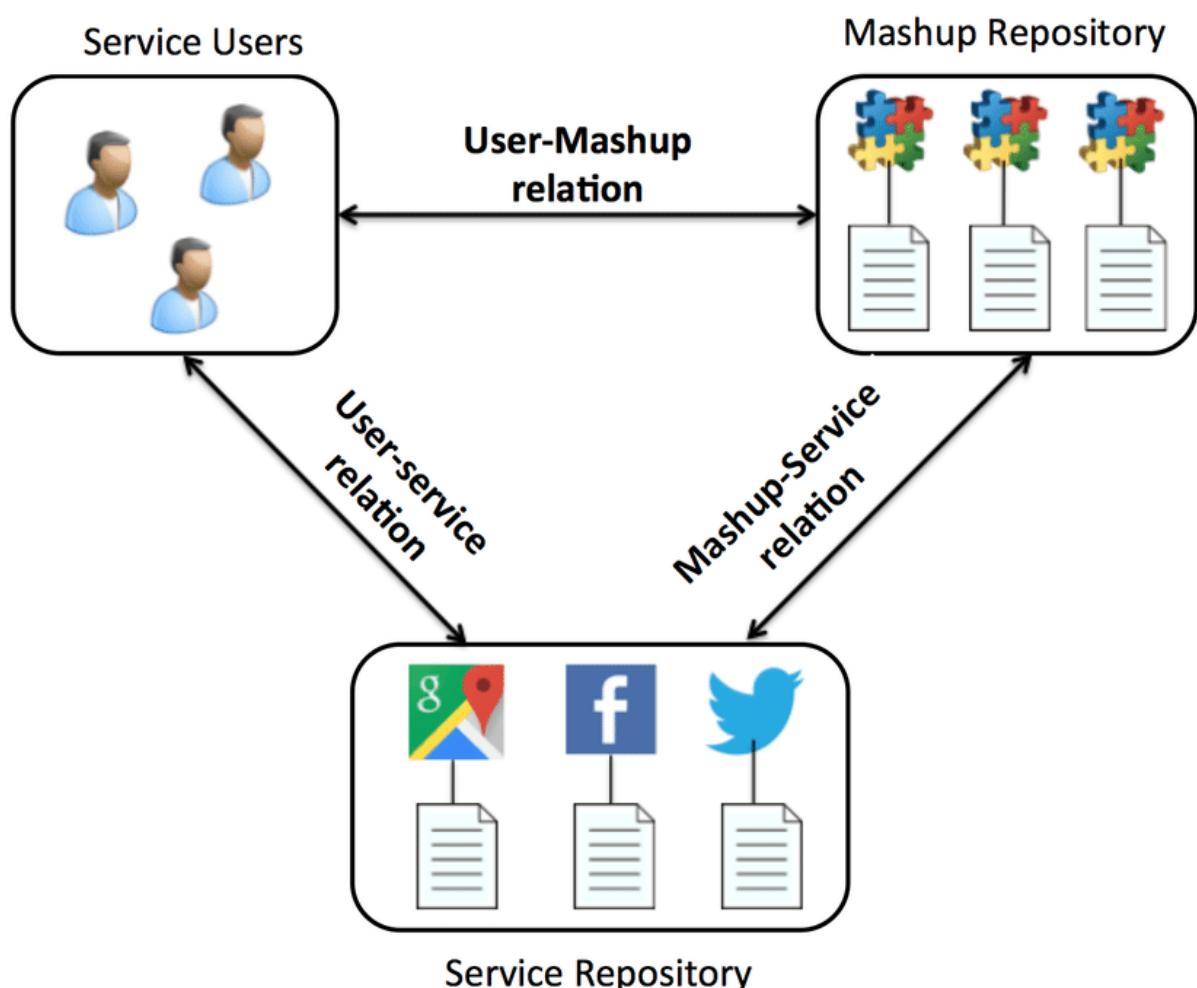
Here's how mashups are relevant in cloud computing:

- 1. Integration of Services:** Cloud platforms often offer a multitude of services and APIs for various purposes such as storage, computing, machine learning, and more. Developers can leverage these services to create mashups that combine functionalities from different services to build innovative applications.
- 2. Customization and Personalization:** Mashups allow developers to tailor solutions to specific needs by combining different cloud services. This enables them to create personalized experiences for users by integrating relevant data and functionalities from disparate sources.

3. Rapid Development: Cloud computing provides the infrastructure and tools necessary for rapid development and deployment of applications. Developers can quickly prototype and deploy mashups by utilizing pre-existing cloud services, thereby accelerating the development process.

4. Scalability and Flexibility: Cloud platforms offer scalability and flexibility, allowing mashup applications to scale seamlessly based on demand. Developers can leverage cloud resources to ensure that their mashup applications can handle varying workloads efficiently.

5. Cost-effectiveness: By utilizing cloud services for building mashups, developers can benefit from a pay-as-you-go model, where they only pay for the resources they consume. This can be cost-effective compared to building and maintaining custom solutions from scratch.



TRUE ENGINEER

A **Mashup** or **hybrid web application** is an application that combines services from several web pages into one to offer a new service. It is common for data integration to be obtained from open **APIs**. The freedom to integrate all these sources is an opportunity to develop new services in a very short time, based on existing functionalities.

Typically mashups rely on using third-party content through public sources or interfaces. Among the most used fonts to create **mashups** are Google Maps, Amazon, Flickr, Youtube, Ebay and **Yahoo**. The ease with which these platforms allow you to integrate your content has led to great success for hybrid web applications.

The term comes from an Anglo-Saxon expression to designate the creation of a song from the mixture of other compositions, usually of very different styles. This origin illustrates very well what a mash up consists of and how it is formed. The term musical has given rise to the name of the hybrid web application from the development of web 2.0.

Advantages of Mashups

Among the main **advantages of mashups** we see the following:

- **Fast content generation.** Integration with existing solutions shortens development times. What used to take weeks of work can now be implemented in just a few minutes.
- **Ease of use.** Anyone with development knowledge can create innovative services by integrating existing functionalities with their own. The design requirements of the same are minimal so creativity is more important than the technical profile.
- **Synergies.** Brands can create innovative experiences without users having to leave their websites. In turn, mashup fonts achieve greater visibility than if their services were only offered through their website.

Services Virtualization Technology

Virtualization

Virtualization is a technology that allows the creation of virtual instances of computing resources, such as servers, storage devices, networks, or operating systems, within a single physical hardware environment. These virtual instances operate independently and can run multiple operating systems and applications simultaneously on the same physical hardware. Virtualization abstracts the underlying hardware, enabling more efficient utilization of resources, better flexibility, scalability, and ease of management in computing environments.

Virtualization is technology that you can use to create virtual representations of servers, storage, networks, and other physical machines. Virtual software mimics the functions of physical hardware to run multiple virtual machines simultaneously on a single physical machine. Businesses use virtualization to use their hardware resources efficiently and get greater returns from their investment. It also powers cloud computing services that help organizations manage infrastructure more efficiently.

Why is virtualization important?

By using virtualization, you can interact with any hardware resource with greater flexibility. Physical servers consume electricity, take up storage space, and need maintenance. You are often limited by physical proximity and network design if you want to access them. Virtualization removes all these limitations by abstracting physical hardware functionality into software. You can manage, maintain, and use your hardware infrastructure like an application on the web.

Virtualization example

Consider a company that needs servers for three functions:

1. Store business email securely
2. Run a customer-facing application
3. Run internal business applications

Each of these functions has different configuration requirements:

- The email application requires more storage capacity and a Windows operating system.
- The customer-facing application requires a Linux operating system and high processing power to handle large volumes of website traffic.
- The internal business application requires iOS and more internal memory (RAM).

To meet these requirements, the company sets up three different dedicated physical servers for each application. The company must make a high initial investment and perform ongoing maintenance and upgrades for one machine at a time. The company also cannot optimize its computing capacity. It pays 100% of the servers' maintenance costs but uses only a fraction of their storage and processing capacities.

To properly understand Kernel-based Virtual Machine (KVM), you first need to understand some basic concepts in *virtualization*. Virtualization is a process that allows a computer to share its hardware resources with multiple digitally separated environments. Each virtualized environment runs within its allocated resources, such as memory, processing power, and storage. With virtualization, organizations can switch between different operating systems on the same server without rebooting.

Virtual machines and hypervisors are two important concepts in virtualization.

Virtual machine

A *virtual machine* is a software-defined computer that runs on a physical computer with a separate operating system and computing resources. The physical computer is called the *host machine* and virtual machines are *guest machines*. Multiple virtual machines can run on a single physical machine. Virtual machines are abstracted from the computer hardware by a hypervisor.

Hypervisor

The *hypervisor* is a software component that manages multiple virtual machines in a computer. It ensures that each virtual machine gets the allocated resources and does not interfere with the operation of other virtual machines. There are two types of hypervisors.

Type 1 hypervisor

A type 1 hypervisor, or bare-metal hypervisor, is a hypervisor program installed directly on the computer's hardware instead of the operating system. Therefore, type 1 hypervisors have better performance and are commonly used by enterprise applications. KVM uses the type 1 hypervisor to host multiple virtual machines on the Linux operating system.

Type 2 hypervisor

Also known as a hosted hypervisor, the type 2 hypervisor is installed on an operating system. Type 2 hypervisors are suitable for end-user computing.

What are the benefits of virtualization?

Virtualization provides several benefits to any organization:

Efficient resource use

Virtualization improves hardware resources used in your data center. For example, instead of running one server on one computer system, you can create a virtual server pool on the same computer system by using and returning servers to the pool as required. Having fewer underlying physical servers frees up space in your data center and saves money on electricity, generators, and cooling appliances.

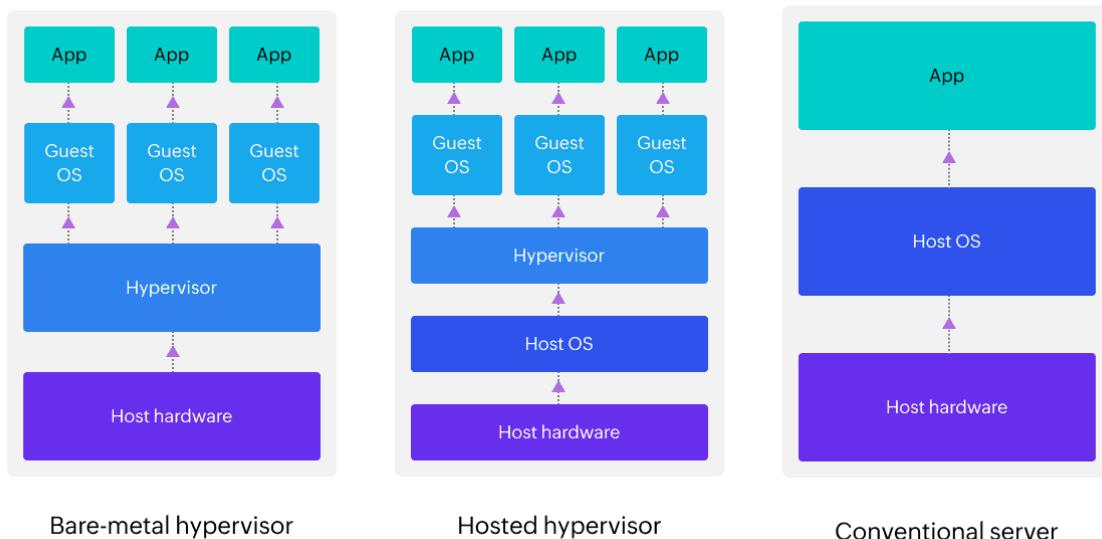
Automated IT management

Now that physical computers are virtual, you can manage them by using software tools. Administrators create deployment and configuration programs to define virtual machine templates. You can duplicate your infrastructure repeatedly and consistently and avoid error-prone manual configurations.

Faster disaster recovery

When events such as natural disasters or cyberattacks negatively affect business operations, regaining access to IT infrastructure and replacing or fixing a physical server can take hours or even days. By contrast, the process takes minutes with virtualized environments. This prompt response significantly improves resiliency and facilitates [business continuity](#) so that operations can continue as scheduled.

Virtualization architecture



Cloud Computing vs Virtualization

CLOUD COMPUTING	VIRTUALIZATION
Cloud computing provides pools and resources which are automated that can be accessed on-demand.	Virtualization is used to make simulated environments through a physical hardware system.
Set-up can be tedious, complicated and a longer process	The Set-up is much simpler when compared to cloud computing
The total operational costs are higher	The operational costs are lower than cloud computing
Cloud computing will provide unlimited storage space	The storage space in virtualization depends on physical server capacity and is limited to its capacity.
Cloud computing requires many dedicated hardware components	A single dedicated hardware can do a great job in virtualization.

Pitfalls of Virtualization

(Pitfall means Disadvantage's)

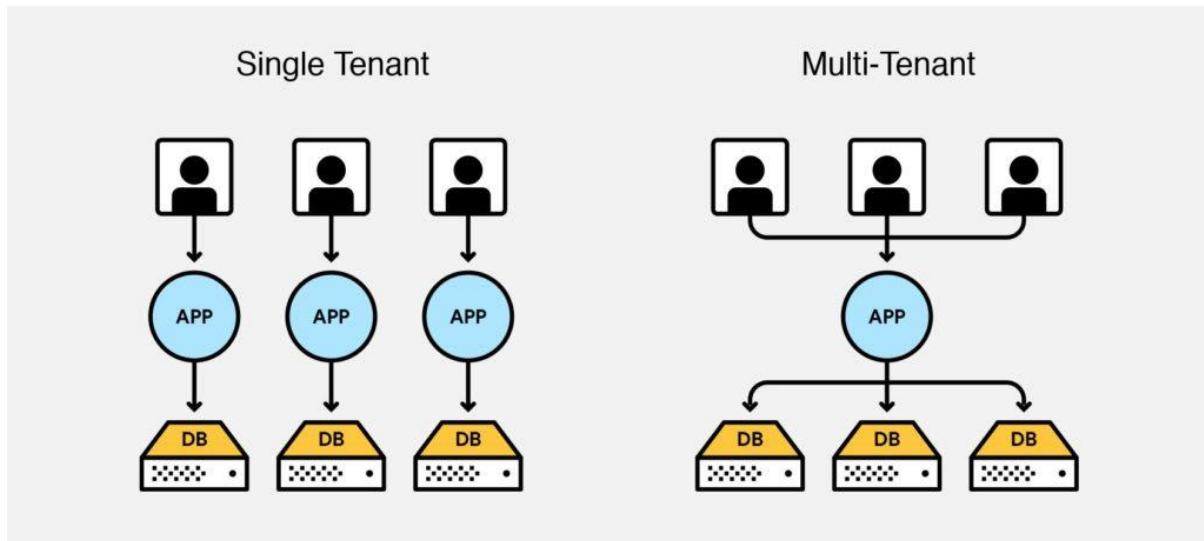
- 1. Performance Overhead:** Virtualization introduces a layer of abstraction between the virtual instances and the physical hardware, which can lead to some performance overhead. This overhead may impact the performance of applications, particularly those that require high computational or I/O resources.
- 2. Resource Contention:** Running multiple virtual instances on the same physical hardware can lead to resource contention, where virtual machines compete for CPU, memory, storage, and network resources. This contention can result in performance degradation or unpredictable behavior, especially during peak usage periods.

TRUE ENGINEER

3. **Security Concerns:** Virtualization introduces new attack vectors and security risks. Vulnerabilities in the hypervisor or misconfigurations of virtual machines can potentially lead to security breaches or unauthorized access to sensitive data. Additionally, co-located virtual instances on the same physical hardware pose a risk of side-channel attacks.
4. **Complexity and Management Overhead:** Managing virtualized environments can be complex, especially as the number of virtual instances grows. Administrators need to deal with tasks such as provisioning, monitoring, patching, and maintaining the virtual infrastructure, which can result in increased management overhead and complexity.
5. **Vendor Lock-in:** Adopting a specific virtualization platform may lead to vendor lock-in, where organizations become dependent on proprietary technologies and may face challenges if they decide to migrate to alternative solutions in the future. This can limit flexibility and increase long-term costs.
6. **Licensing and Compliance:** Virtualization can have implications for software licensing and compliance. Organizations need to ensure compliance with software licensing agreements, which may have specific requirements or restrictions related to virtualized environments. Failure to comply can result in legal and financial consequences.
7. **Single Point of Failure:** While virtualization can enhance flexibility and scalability, it also introduces a single point of failure in the form of the hypervisor. If the hypervisor fails, it can potentially impact all virtual instances running on the physical hardware, leading to service disruptions or downtime.

Multitenant software

Multitenant software is an application architecture designed to serve multiple customers or tenants (such as organizations or individual users) from a single instance of the software. In a multitenant architecture, each tenant's data and configuration are logically isolated from one another, allowing them to share the same underlying infrastructure and resources while maintaining data privacy, security, and customization.



Multi-tenant architecture, more commonly referred to as multi-tenancy, is a software architecture where multiple instances of an application run on the same physical server. The same server is then responsible for serving multiple tenants simultaneously.

This type of build allows companies to allocate a single infrastructure to several end users, rather than individually managing the maintenance and updates of multiple environments.

Here are some key aspects and characteristics of multitenant software:

- 1. Shared Infrastructure:** Multitenant software shares the same physical or virtual infrastructure, including servers, storage, and networking resources, among multiple tenants. This allows for efficient resource utilization and cost savings compared to deploying separate instances for each tenant.
- 2. Data Isolation:** Despite sharing the same infrastructure, each tenant's data is logically isolated from other tenants. This ensures that each tenant has exclusive access to their data and prevents unauthorized access or data leakage between tenants.
- 3. Customization and Configuration:** Multitenant software often provides mechanisms for tenants to customize and configure their environments according to their specific requirements. This may include customizing user interfaces, workflows, access controls, and branding to suit each tenant's needs.
- 4. Scalability:** Multitenant architectures are designed to scale horizontally to accommodate the needs of multiple tenants. As the number of tenants or the workload of individual tenants increases, the software can dynamically allocate resources to meet demand while maintaining performance and responsiveness.

TRUE ENGINEER

5. **Upgrade and Maintenance:** Upgrades and maintenance tasks are typically performed centrally by the software provider, ensuring that all tenants benefit from the latest features, bug fixes, and security patches without requiring individual tenant intervention. This reduces the administrative overhead for tenants and ensures consistency across the environment.
6. **Subscription-based Pricing:** Multitenant software often employs a subscription-based pricing model, where tenants pay for the services they use on a recurring basis. This model allows for flexible billing based on usage metrics or predefined tiers and provides predictable costs for tenants.
7. **Security and Compliance:** Multitenant software must implement robust security measures to ensure the privacy and integrity of each tenant's data. This includes encryption, access controls, authentication mechanisms, and compliance with relevant regulations and industry standards.

Multi-tenancy using cloud data stores



Multi-tenancy using cloud data stores refers to a software architecture where multiple tenants share a single instance of a cloud-based data storage system. In this setup, each tenant's data is logically isolated from other tenants, allowing them to operate independently while sharing the same underlying infrastructure.

TRUE ENGINEER

- 1. Shared Infrastructure:** Cloud data stores provide a centralized platform for storing and managing data, offering scalability, reliability, and accessibility over the internet. In a multi-tenant setup, multiple tenants share the same cloud-based data store, leveraging the cloud provider's infrastructure.
- 2. Logical Isolation:** While tenants share the same data store, their data is logically isolated from one another to ensure privacy, security, and data integrity. This isolation is typically achieved through data partitioning, where each tenant's data is stored in separate partitions or namespaces within the data store.
- 3. Tenant Identification:** Multi-tenancy in cloud data stores often involves mechanisms for identifying and authenticating tenants, ensuring that each tenant has appropriate access permissions to their own data while preventing unauthorized access to other tenants' data. This may involve using tenant-specific identifiers, access tokens, or authentication keys.
- 4. Customization and Configuration:** Cloud data stores may offer features for customizing and configuring data management settings for each tenant, such as defining access controls, data retention policies, backup schedules, and data encryption requirements. This allows tenants to tailor the data storage environment to their specific needs and compliance requirements.
- 5. Scalability and Performance:** Multi-tenancy in cloud data stores enables efficient resource utilization and scalability, as tenants share the same infrastructure resources while maintaining isolation between their data. The data store can dynamically allocate resources to handle varying workloads and scale horizontally to accommodate growth in the number of tenants or data volume.
- 6. Cost-effectiveness:** Sharing a single instance of a cloud data store among multiple tenants can result in cost savings compared to maintaining separate data stores for each tenant. Tenants benefit from economies of scale and pay only for the storage and resources they consume, typically following a pay-as-you-go pricing model.
- 7. Maintenance and Updates:** Cloud data store providers are responsible for managing and maintaining the underlying infrastructure, including hardware maintenance, software updates, security patches, and backups. This relieves tenants from the burden of managing infrastructure maintenance tasks and ensures that the data store remains secure and up-to-date.

What are the benefits of multitenancy?

Many of the benefits of cloud computing are only possible because of multitenancy. Here are two crucial ways multitenancy improves cloud computing:

Better use of resources: One machine reserved for one tenant is not efficient, as that one tenant is not likely to use all of the machine's computing power. By sharing machines among multiple tenants, use of available resources is maximized.

Lower costs: With multiple customers sharing resources, a cloud vendor can offer their services to many customers at a much lower cost than if each customer required their own dedicated infrastructure.

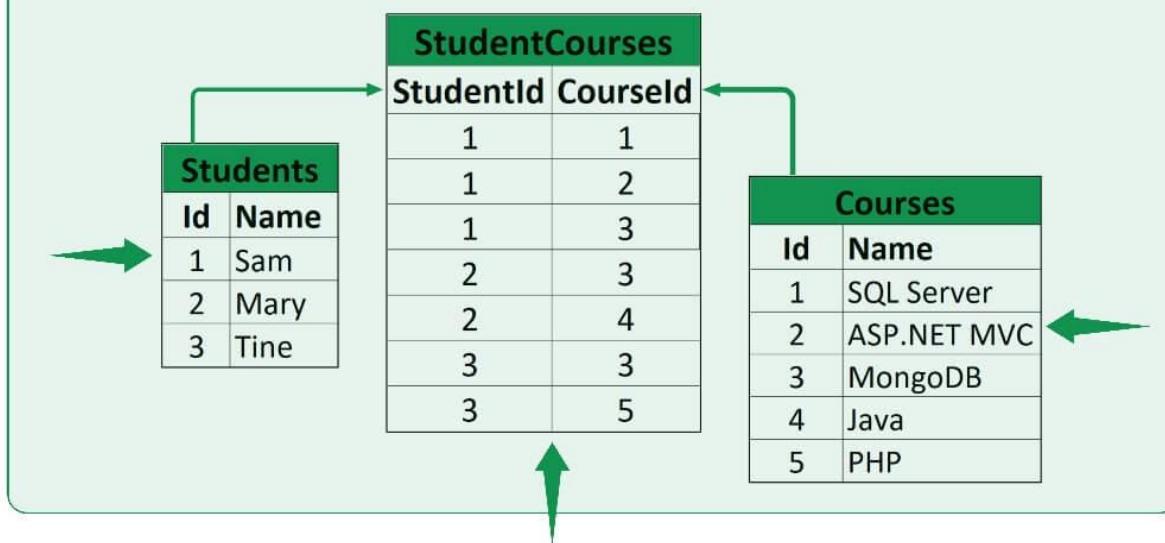
Unit 03

Relational databases

A relational database is a type of database that organizes data into tables consisting of rows and columns, with each row representing a record and each column representing a specific attribute or field of that record. The structure of a relational database is defined by a schema, which specifies the tables, fields, relationships, and constraints within the database.

- 1. Scalability:** Cloud-based relational databases can scale up or down easily based on demand. Cloud providers offer various scaling options, such as vertical scaling (increasing the resources of a single instance) or horizontal scaling (adding more instances to distribute the load).
- 2. Flexibility:** Cloud-based relational databases provide flexibility in terms of deployment options. You can choose from fully managed database services provided by cloud providers (e.g., Amazon RDS, Google Cloud SQL, Azure SQL Database) or deploy your own database on virtual machines or containers within the cloud environment.
- 3. High Availability:** Cloud providers offer built-in high availability features for relational databases, such as automated backups, failover mechanisms, and multi-region replication. This ensures that your database remains accessible and resilient to failures.
- 4. Security:** Cloud providers implement robust security measures to protect data in transit and at rest. They offer features like encryption, identity and access management, network isolation, and compliance certifications to ensure data security and regulatory compliance.
- 5. Cost Efficiency:** Cloud-based relational databases often follow a pay-as-you-go pricing model, allowing you to pay only for the resources you use. Additionally, cloud providers offer cost optimization tools and features to help you manage and optimize your database costs.
- 6. Managed Services:** Fully managed database services provided by cloud providers abstract away the complexity of database management tasks such as provisioning, patching, backups, and monitoring. This allows you to focus more on developing applications rather than managing infrastructure.
- 7. Integration with Other Cloud Services:** Cloud-based relational databases seamlessly integrate with other cloud services such as compute, storage, analytics, and machine learning services. This enables you to build comprehensive and scalable cloud-native applications.

Relational Database



Key characteristics of relational databases include:

- 1. Tabular Structure:** Data is organized into tables, with each table consisting of rows and columns. Each row represents a unique record, and each column represents a specific attribute of that record.
- 2. Relationships:** Relational databases support relationships between tables, allowing data to be linked and queried across multiple tables. Common types of relationships include one-to-one, one-to-many, and many-to-many.
- 3. ACID Transactions:** Relational databases ensure data integrity and consistency through the use of ACID (Atomicity, Consistency, Isolation, Durability) transactions. Transactions ensure that database operations are executed reliably and that the database remains in a consistent state even in the event of failures.
- 4. SQL (Structured Query Language):** Relational databases use SQL as the standard language for querying and manipulating data. SQL provides a powerful and standardized syntax for performing operations such as SELECT (querying), INSERT (inserting), UPDATE (updating), and DELETE (deleting) data.
- 5. Normalization:** Relational databases are designed to adhere to the principles of normalization, which involves organizing data to minimize redundancy and dependency. Normalization helps improve data integrity and reduce storage space.

Cloud file systems

Cloud File Storage

Cloud file storage is a method for storing data in the cloud that provides servers and applications access to data through shared file systems. This compatibility makes cloud file storage ideal for workloads that rely on shared file systems and provides simple integration without code changes.

Cloud File System

A cloud file system is a hierarchical storage system in the cloud that provides shared access to file data. Users can create, delete, modify, read, and write files, as well as organize them logically in directory trees for intuitive access.

Cloud file systems are storage systems designed to store and manage data in cloud computing environments. They offer scalable, distributed, and highly available storage solutions for cloud-based applications and services. Here are some key aspects of cloud file systems:

1. **Scalability:** Cloud file systems are designed to scale horizontally, allowing them to handle large volumes of data and increasing performance as the workload grows. They can dynamically allocate resources to accommodate changing storage needs without requiring manual intervention.
2. **Distributed Architecture:** Cloud file systems distribute data across multiple nodes or servers within a cloud environment. This distributed architecture improves fault tolerance and availability by replicating data across multiple locations and ensuring redundancy.
3. **High Availability:** Cloud file systems are built to provide high availability, ensuring that data remains accessible even in the event of hardware failures or network issues. They employ techniques such as data replication, automatic failover, and load balancing to maintain service uptime.
4. **Data Consistency:** Cloud file systems maintain data consistency across distributed nodes by implementing synchronization protocols and distributed locking mechanisms. This ensures that all nodes have access to the most up-to-date version of the data and that changes are propagated reliably.

5. Performance Optimization: Cloud file systems optimize performance through caching mechanisms, data compression, and parallel processing techniques. They are designed to deliver low-latency access to data, especially for latency-sensitive applications such as real-time analytics and media streaming.

6. Security: Cloud file systems provide robust security features to protect data in transit and at rest. They offer encryption, access control mechanisms, auditing capabilities, and compliance certifications to ensure data privacy and regulatory compliance.

7. Integration with Cloud Services: Cloud file systems seamlessly integrate with other cloud services such as compute, analytics, and data processing services. This allows applications to access and manipulate data stored in the file system using APIs and SDKs provided by the cloud provider.

GFS(Google File System)

GFS is a file system designed to handle batch workloads with lots of data.

The system is **distributed**: multiple machines store copies of every file, and multiple machines try to read/write the same file. GFS was originally designed for Google's use case of searching and indexing the web. So, at its core, GFS addresses the following concerns:

- **Fault Tolerance:** Google uses **commodity machines** because they are cheap and easy to acquire, but the software behind GFS needs to be robust to handle failures of machines, disks, and networks.
- **Large Files:** It's assumed most files are large (i.e. ≥ 100 MB). Small files are supported, but not optimized for.
- **Optimize for Reads + Appends:** The system is optimized for **reading** (specifically large streaming reads) or **appending** because web crawling and indexing heavily relied on these operations.
- **High and Consistent Bandwidth:** It's acceptable to have slow operations now and then, but the overall amount of data flowing through the system should be **consistent**. Again, this stems from Google's crawling and indexing purposes.

GFS Architecture

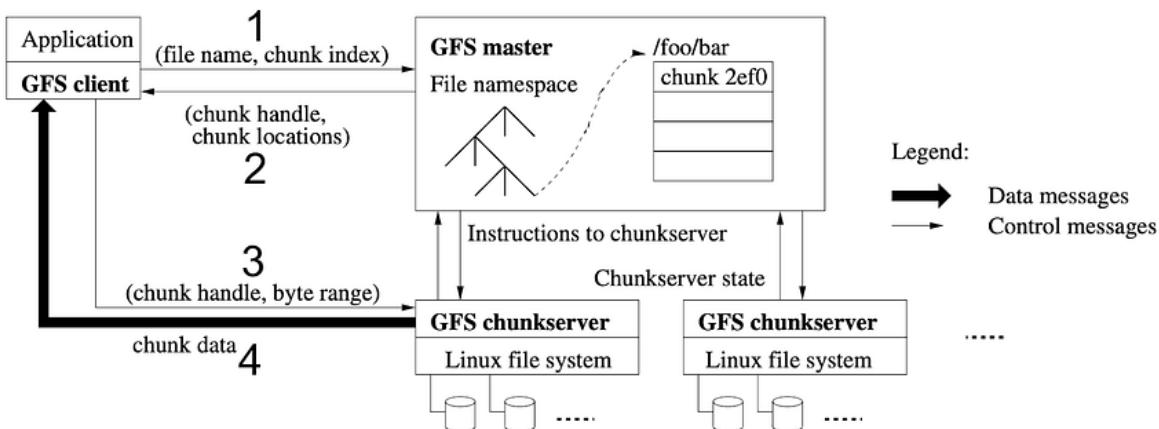


Figure 1: GFS Architecture

- Scalability:** GFS is designed to handle large-scale data storage requirements, spanning thousands of commodity servers across multiple data centers. It can store petabytes of data and scale horizontally by adding more servers to the cluster.
- Fault Tolerance:** GFS is fault-tolerant, meaning it can continue to operate even in the presence of hardware failures, network issues, or other types of failures. It achieves fault tolerance through data replication and automatic recovery mechanisms.
- Distributed Architecture:** GFS distributes data across multiple storage servers called "chunkservers" and manages metadata through a centralized "master" server. This distributed architecture allows GFS to handle large volumes of data and provide high availability and reliability.
- Reliability:** GFS ensures data reliability by replicating data across multiple chunkservers. By default, data is replicated at least three times to ensure redundancy and fault tolerance. If a chunkserver fails, GFS can retrieve the data from other replicas.
- Streaming Reads and Writes:** GFS is optimized for streaming reads and writes, making it suitable for data-intensive applications such as web indexing, data processing, and multimedia streaming. It achieves high throughput by minimizing seek times and maximizing data transfer rates.
- Consistency Model:** GFS provides a relaxed consistency model, allowing for eventual consistency rather than strong consistency. This means that updates to the file system may not be immediately visible to all clients, but consistency is eventually achieved through periodic synchronization.
- Automatic Load Balancing:** GFS automatically balances the load across chunkservers to ensure optimal performance and resource utilization. It redistributes data and workload dynamically based on the current state of the system.
- Integration with MapReduce:** GFS is tightly integrated with Google's MapReduce framework, allowing distributed processing of large datasets stored in the file system. This integration enables parallel execution of data-intensive tasks across multiple nodes in the GFS cluster.

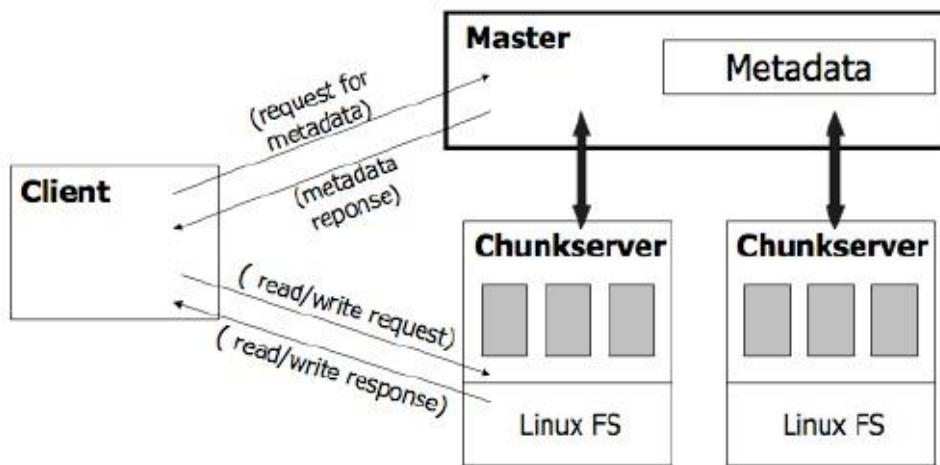


Figure 1

In the Google File System (GFS), the architecture consists of two main components: the master server and the chunk servers.

1. Master Server:

- The master server is responsible for coordinating access to the file system and managing metadata.
- It maintains metadata about files, including their locations, sizes, and access permissions.
- The master server keeps track of which chunk servers store which data chunks and handles operations such as file creation, deletion, and renaming.
- It assigns unique identifiers to each chunk and keeps track of the chunk replicas' locations.
- The master server also monitors the health of chunk servers, detects failures, and initiates replica placement and re-replication as necessary.
- The master server typically runs on a single node in the GFS cluster and is designed to be highly available and fault-tolerant.

2. Chunk Servers:

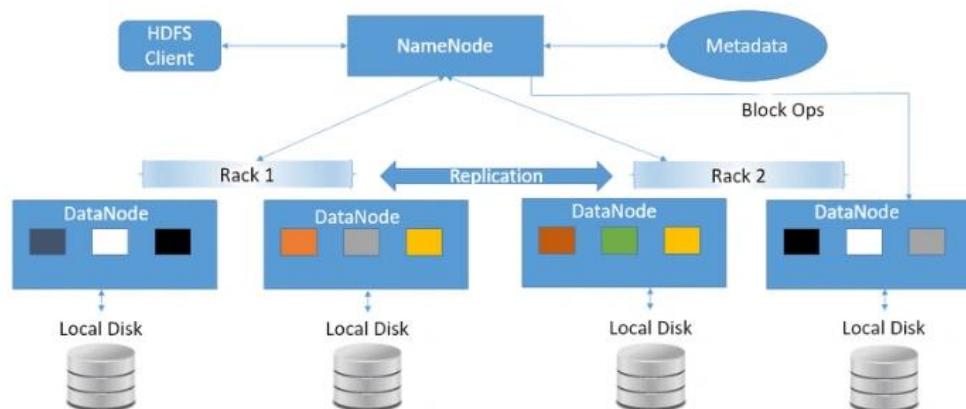
- Chunk servers are responsible for storing and serving data chunks.
- Each chunk server manages multiple data chunks, typically storing multiple replicas of each chunk for fault tolerance.
- Chunk servers handle read and write requests from clients, as well as data replication and recovery operations initiated by the master server.
- They continuously report their status and availability to the master server and respond to commands and instructions sent by the master.
- Chunk servers are designed to be stateless, meaning they do not store any persistent metadata and rely on the master server for metadata management.
- Chunk servers are distributed across multiple nodes in the GFS cluster, providing fault tolerance and scalability.

HDFS (Hadoop Distributed File System)

Hadoop Distributed File System (HDFS) is a distributed file system designed to store and manage large volumes of data across a cluster of commodity hardware. It is a core component of the Apache Hadoop framework and is widely used for storing data in big data analytics and processing applications. Here are some key features of HDFS:

- 1. Distributed Storage:** HDFS distributes data across multiple servers or nodes in a cluster, allowing it to store large datasets that exceed the capacity of a single server. Data is divided into blocks, typically 128 MB or 256 MB in size, and distributed across the cluster.
- 2. Replication:** HDFS achieves fault tolerance and data reliability through replication. Each data block is replicated across multiple nodes in the cluster, typically three replicas by default. If a node fails or becomes unavailable, HDFS can retrieve the data from one of the replicas stored on other nodes.
- 3. Streaming Data Access:** HDFS is optimized for streaming data access rather than random access. It is well-suited for applications that require large sequential reads and writes, such as batch processing, data ingestion, and analytics.
- 4. Scalability:** HDFS is highly scalable and can scale out by adding more nodes to the cluster. It can handle petabytes or even exabytes of data by distributing the storage and processing load across multiple nodes.
- 5. Fault Tolerance:** HDFS is designed to be fault-tolerant, meaning it can continue to operate even in the presence of node failures, network issues, or other types of failures. It achieves fault tolerance through data replication, metadata redundancy, and automatic recovery mechanisms.
- 6. Single NameSpace:** HDFS provides a unified namespace for all files and directories stored in the cluster, similar to a traditional file system. This allows users to interact with the file system using familiar commands and APIs.
- 7. Data Integrity:** HDFS ensures data integrity through checksums, which are computed for each data block and verified during read operations. If a checksum mismatch is detected, HDFS can retrieve the data from another replica to ensure data consistency.
- 8. Integration with Hadoop Ecosystem:** HDFS integrates seamlessly with other components of the Hadoop ecosystem, such as MapReduce, Apache Hive, Apache Spark, and Apache HBase. This integration enables distributed processing of large datasets stored in HDFS using various data processing frameworks and tools.

HDFS Architecture

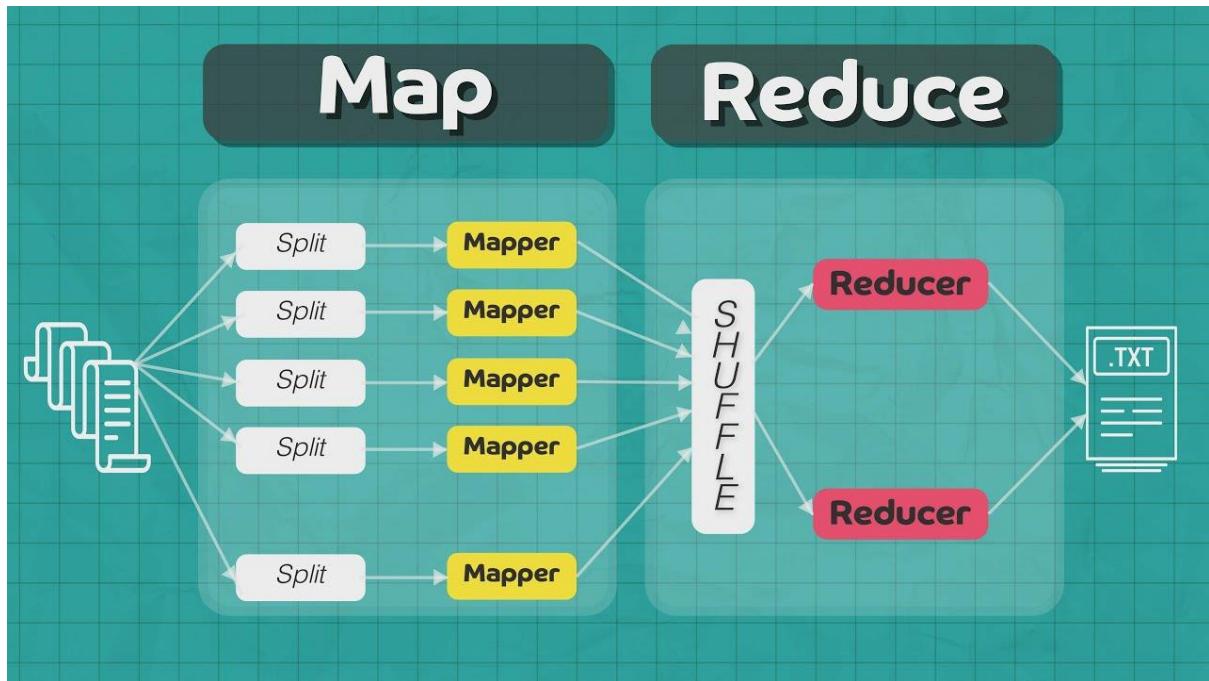


GFS VS HDFS

GFS VS. HDFS

GFS	HDFS
Master	NameNode
ChunkServer	DataNode
Operation Log	Journal, Edit Log
Chunk	Block
Random file writes possible	Only append is possible
Multiple writer/reader model	Single writer/multiple reader model
Default chunk size: 64MB	Default block size: 128MB

Map Reduce



Map Reduce Working

1. Map Phase:

- In the Map phase, data processing tasks are divided into smaller sub-tasks called "map tasks".
- Each map task processes a portion of the input data independently and in parallel.
- The input data is typically stored in a distributed file system, such as Hadoop Distributed File System (HDFS).
- The user-defined map function is applied to each input record, generating a set of intermediate key-value pairs.

2. Shuffle and Sort Phase:

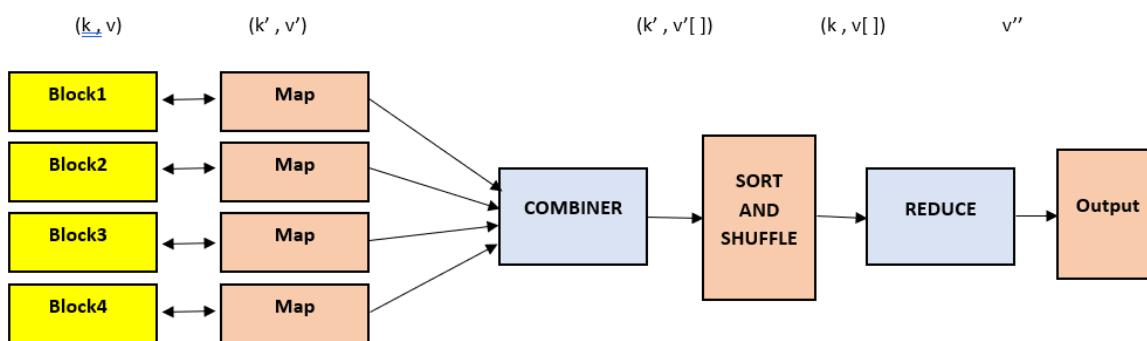
- In the Shuffle and Sort phase, the intermediate key-value pairs produced by the map tasks are shuffled and grouped by key.
- Key-value pairs with the same key are grouped together and sorted by key to prepare for the subsequent Reduce phase.
- This phase ensures that all values associated with the same key are processed together by the same reduce task.

3. Reduce Phase:

- In the Reduce phase, the grouped and sorted intermediate key-value pairs are processed by user-defined reduce functions.
- Each reduce task processes a subset of the intermediate key-value pairs, typically grouped by key.
- The user-defined reduce function is applied to each group of values with the same key, producing the final output key-value pairs.

4. Output:

- The output of the Reduce phase is the final result of the MapReduce job, consisting of key-value pairs generated by the reduce tasks.
- The output can be stored in a distributed file system or used as input for subsequent MapReduce jobs or other data processing tasks.



Benefits of Map Reduce

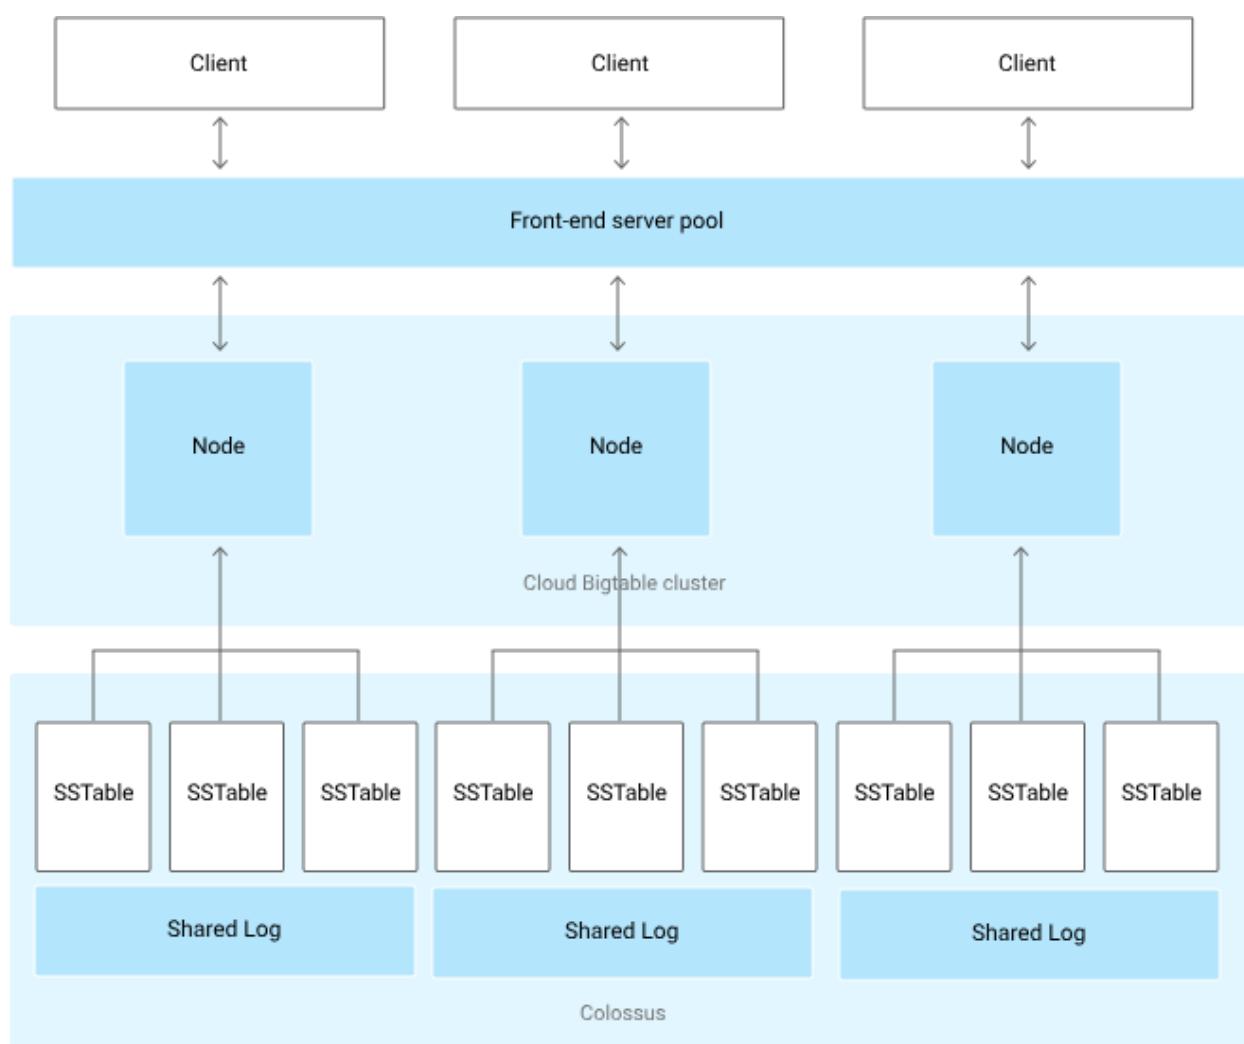
- **Scalability:** MapReduce scales horizontally by distributing data processing tasks across multiple nodes in a cluster, allowing it to handle large-scale datasets and parallelize computation.
- **Fault Tolerance:** MapReduce is designed to be fault-tolerant, automatically handling node failures and task re-execution to ensure reliable and consistent results.
- **Data Locality:** MapReduce leverages data locality by moving computation closer to the data, minimizing data transfer over the network and optimizing performance.
- **Simplicity:** MapReduce provides a simple and easy-to-understand programming model, allowing developers to focus on the logic of their map and reduce functions without worrying about low-level details of distributed computing.

Big Table

In cloud computing, a Big Table is a term that typically refers to a distributed storage system for managing structured data. It's designed to handle massive amounts of data across a distributed infrastructure, providing scalability, fault tolerance, and high performance for applications that require storage and retrieval of large datasets.

Google's Bigtable is one of the most well-known implementations of a Big Table in cloud computing. It's a NoSQL, wide-column database service designed to handle massive amounts of structured data and is a core component of Google Cloud Platform's (GCP) storage offerings.

Bigtable is optimized for storing large amounts of semi-structured data, such as web indexing data, time-series data, or analytics data. It's known for its scalability, allowing users to seamlessly scale their storage and throughput as their data needs grow.



Key Points of Big Table

1. Large-Scale Storage

BigTable offers you the potential for capturing high amounts of data and use it for machine learning and analytics.

2. Integration with ML Tools

With BigTable, you can run Machine Learning algorithms for deriving recommendations and predictions.

3. High Throughput

All the data that intends to change rapidly can be accommodated with a high count of reading and write throughputs.

4. Latency is Low

Low latency resembles high-speed on-site performance. You will get the response for lookups or data searches in a matter of milliseconds.

5. High Availability

BigTable can constantly operate without downtime or breakdown to serve the customers without fail.

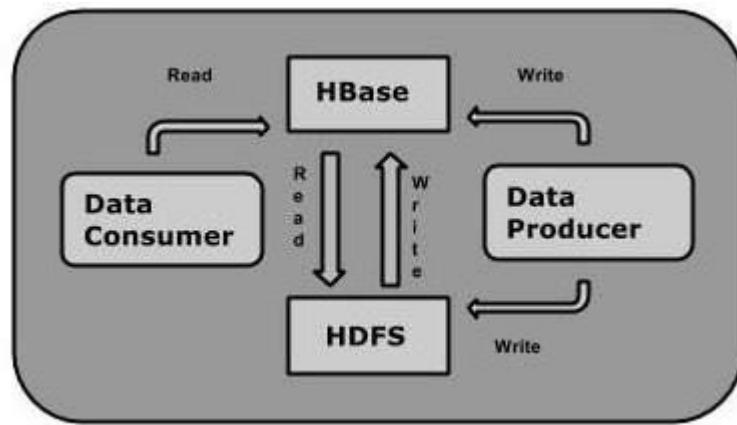
6. High Scalability

BigTable is scalable to cater to many requests and users automatically without any overloading aspect.

7. Fully-Managed Database

BigTable, being fully managed, allows the developers to direct their focus upon app development aspects instead of giving their attention to databases. Google BigTable is handling that part seamlessly!

H Base

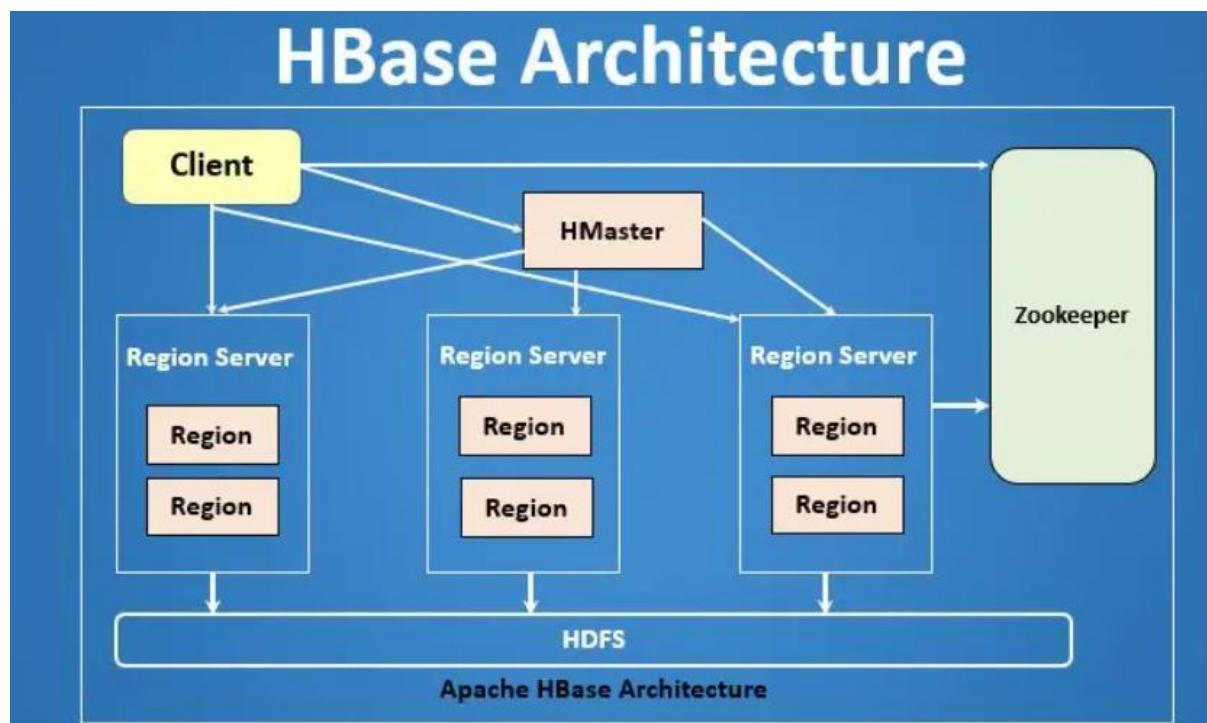


HBase is another popular distributed, scalable, and NoSQL database system, often used in the context of big data and cloud computing environments. It's built on top of the Hadoop Distributed File System (HDFS) and is part of the Apache Hadoop project. HBase is modeled after Google's Bigtable, and it shares many similarities with Bigtable in terms of its architecture and use cases.

Here are some key aspects of HBase:

- 1. Column-family based:** Like Bigtable, HBase organizes data into tables, rows, and column families. Each row in an HBase table can have a flexible number of columns, grouped into column families. This allows for efficient storage of sparse data and provides flexibility in schema design.
- 2. Distributed and scalable:** HBase is designed to scale horizontally across a cluster of machines. It automatically partitions and distributes data across multiple nodes in the cluster, allowing it to handle large datasets spanning hundreds or thousands of nodes.
- 3. Highly available:** HBase provides replication and failover mechanisms to ensure high availability of data. Data is automatically replicated across multiple nodes, and in the event of node failures, HBase can transparently redirect requests to replicas, minimizing downtime.

4. **Consistent and strong consistency:** HBase provides strong consistency guarantees for read and write operations. It uses techniques such as distributed consensus and versioning to ensure that reads and writes are consistent across the cluster.
5. **Integration with Hadoop ecosystem:** HBase is tightly integrated with other components of the Hadoop ecosystem, such as HDFS, MapReduce, and Apache Spark. This makes it easy to build data processing pipelines that leverage HBase for storing and accessing large-scale data.
6. **Use cases:** HBase is well-suited for use cases that require real-time random read/write access to large datasets, such as real-time analytics, social media platforms, recommendation systems, and Internet of Things (IoT) applications.

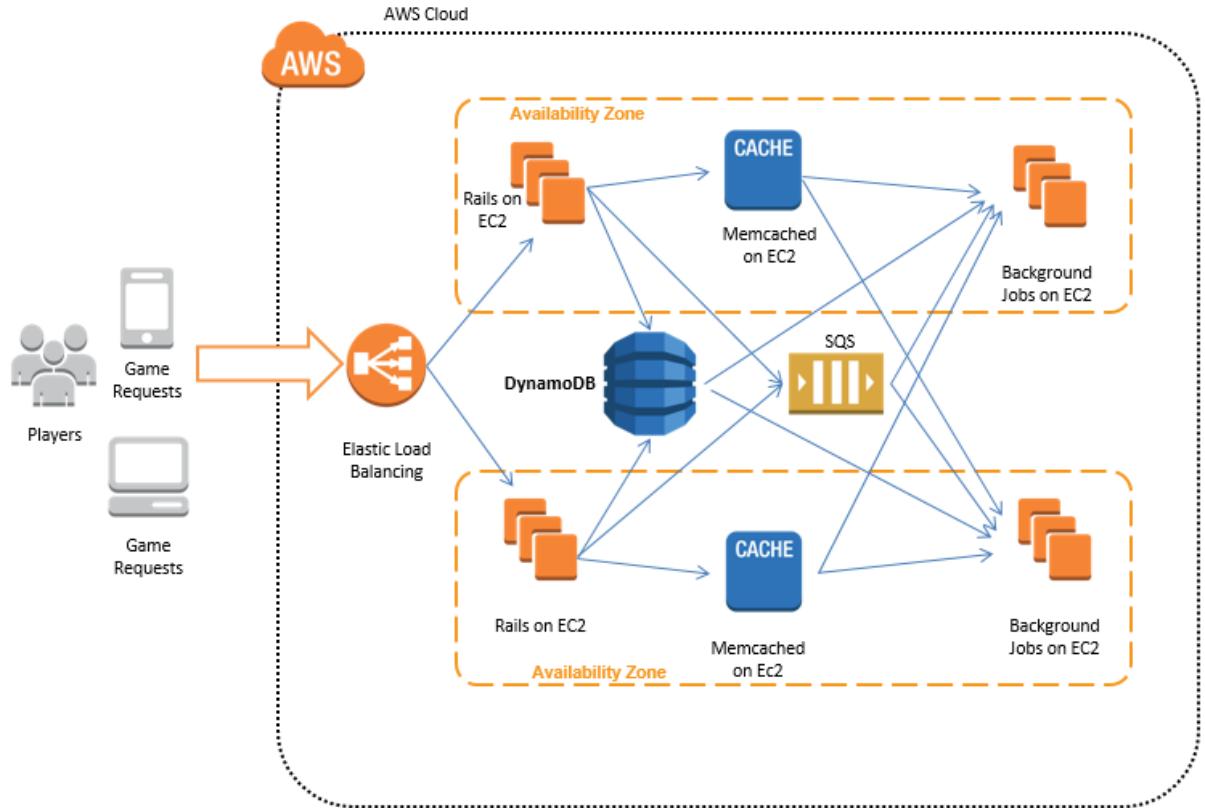


Dynamo

Dynamo is a highly available and scalable distributed data store developed by Amazon.com for managing key-value data. It was designed to address the need for a database system that could handle the massive scale and high availability requirements of Amazon's e-commerce platform. Dynamo influenced the development of several subsequent NoSQL databases and is considered a pioneering system in the field of distributed databases.

Here are some key characteristics of Dynamo:

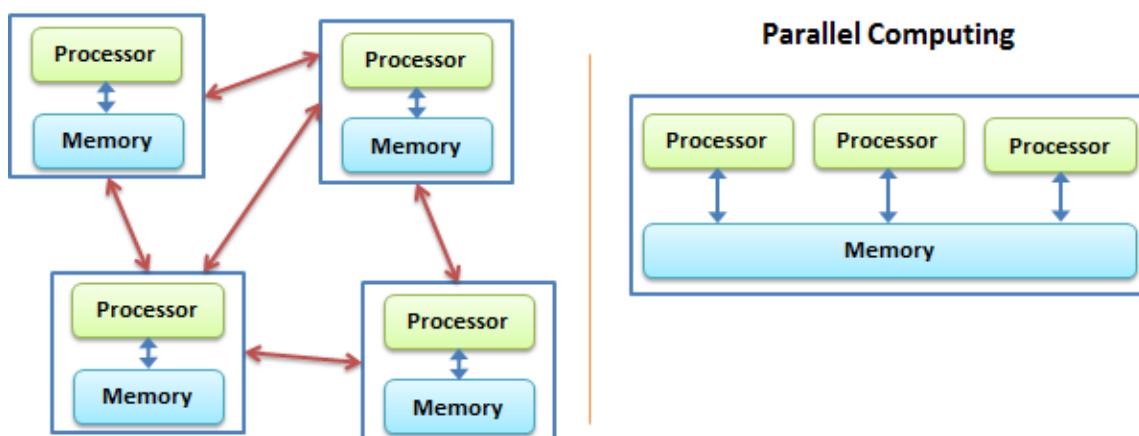
1. **Key-Value Store:** Dynamo is a key-value store, meaning that data is stored and retrieved using keys. Each key is associated with a value, which can be any type of data, such as a string, binary object, or JSON document.
2. **Partitioning and Replication:** Dynamo partitions data across multiple nodes in a cluster to achieve scalability. Each node in the cluster is responsible for a range of keys, and the system uses consistent hashing to determine which node is responsible for each key. Data is also replicated across multiple nodes to provide fault tolerance and high availability.
3. **Consistency and Availability Trade-off:** Dynamo employs a configurable consistency model that allows users to choose between strong consistency and eventual consistency. Strong consistency ensures that all read and write operations are immediately reflected across all replicas, while eventual consistency allows for faster reads by relaxing consistency requirements and asynchronously propagating updates.
4. **Incremental Scalability:** Dynamo is designed to scale incrementally by adding more nodes to the cluster as the workload or data size grows. The system automatically rebalances data and redistributes load across nodes to accommodate changes in the cluster size.
5. **Conflict Resolution:** In the event of conflicting updates to the same key, Dynamo provides mechanisms for conflict resolution. Typically, conflicts are resolved based on application-specific logic or using timestamp-based reconciliation.
6. **Integration with Amazon Web Services (AWS):** Dynamo is tightly integrated with AWS and is available as a fully managed service called Amazon DynamoDB. DynamoDB provides a scalable, low-latency, and fully managed NoSQL database solution that is suitable for a wide range of applications.



Parallel Computing

Parallel computing refers to the simultaneous execution of multiple computational tasks, with the goal of speeding up the overall computation by dividing the workload among multiple processing units, such as CPU cores, GPUs, or even distributed computing nodes. This approach contrasts with serial computing, where tasks are executed sequentially on a single processing unit.

Distributed Computing



TRUE ENGINEER

1. **Task Parallelism:** In task parallelism, different tasks or processes are executed concurrently. Each task may perform different operations on different sets of data. This approach is often used in applications where the tasks are independent and can be executed in parallel without needing to communicate or synchronize with each other extensively.
2. **Data Parallelism:** Data parallelism involves dividing a single task into smaller sub-tasks, each of which operates on different subsets of data. These sub-tasks are then executed concurrently on multiple processing units. Data parallelism is common in applications that involve processing large datasets, such as image processing, numerical simulations, and machine learning algorithms.
3. **Parallel Architectures:** Parallel computing can be implemented using various architectures, including shared-memory systems, distributed-memory systems, and hybrid architectures that combine elements of both. Shared-memory systems allow multiple processing units to access the same memory, while distributed-memory systems have separate memory spaces for each processing unit.
4. **Parallel Programming Models:** Parallel programming models provide abstractions and constructs for expressing parallelism in software. Common parallel programming models include message passing (e.g., MPI), shared-memory multiprocessing (e.g., OpenMP), and dataflow-based models (e.g., Apache Spark).
5. **Synchronization and Communication:** In parallel computing, synchronization and communication are crucial for coordinating the execution of tasks and ensuring correctness. Techniques such as locks, barriers, and message passing are used to synchronize access to shared resources and exchange data between parallel tasks.
6. **Scalability:** Scalability refers to the ability of a parallel computing system to efficiently utilize additional resources as the problem size or workload increases. Scalability is essential for handling large-scale computations and accommodating growth in data volume or processing requirements.

Enterprise batch processing

Enterprise batch processing refers to the automated execution of a series of computational tasks or jobs on a large scale within an organization's computing environment. These tasks are typically scheduled to run at specific times or triggered by certain events, such as the availability of new data or the completion of previous jobs. Batch processing is commonly used in enterprise environments for a variety of purposes, including data processing, report generation, ETL (Extract, Transform, Load) operations, and system maintenance.

Unit 04

Cloud Security

Cloud security refers to the set of technologies, policies, controls, and procedures designed to protect data, applications, and infrastructure in cloud computing environments. As businesses increasingly migrate their operations to the cloud, ensuring the security of cloud-based resources becomes paramount. Here are some key aspects of cloud security:

- 1. Data Encryption:** Encrypting data both in transit and at rest helps prevent unauthorized access. Encryption keys should be carefully managed to ensure only authorized users can decrypt the data.
- 2. Identity and Access Management (IAM):** Implementing robust IAM controls ensures that only authenticated and authorized individuals can access cloud resources. This involves practices like multi-factor authentication (MFA), role-based access control (RBAC), and least privilege principles.
- 3. Network Security:** Securely configuring virtual networks, firewalls, and access controls within cloud environments helps prevent unauthorized access and protects against network-based attacks.
- 4. Vulnerability Management:** Regularly scanning cloud environments for vulnerabilities and promptly patching or remediating any identified issues helps reduce the risk of exploitation by attackers.
- 5. Security Monitoring and Logging:** Continuous monitoring of cloud infrastructure and applications, along with thorough logging of events, enables rapid detection and response to security incidents.
- 6. Compliance and Governance:** Ensuring compliance with relevant regulations and industry standards, as well as implementing robust governance practices, helps mitigate risks and maintain trust with customers and regulators.
- 7. Incident Response:** Developing and regularly testing incident response plans helps organizations effectively respond to security breaches or incidents in cloud environments.

TRUE ENGINEER

8. Provider Security Measures: Understanding the security measures provided by the cloud service provider (CSP) is crucial. Many CSPs offer security features and compliance certifications, but it's important for organizations to understand their own responsibilities for security within the shared responsibility model.

9. Data Loss Prevention (DLP): Implementing DLP policies and controls helps prevent sensitive data from being lost, stolen, or exposed within cloud environments.

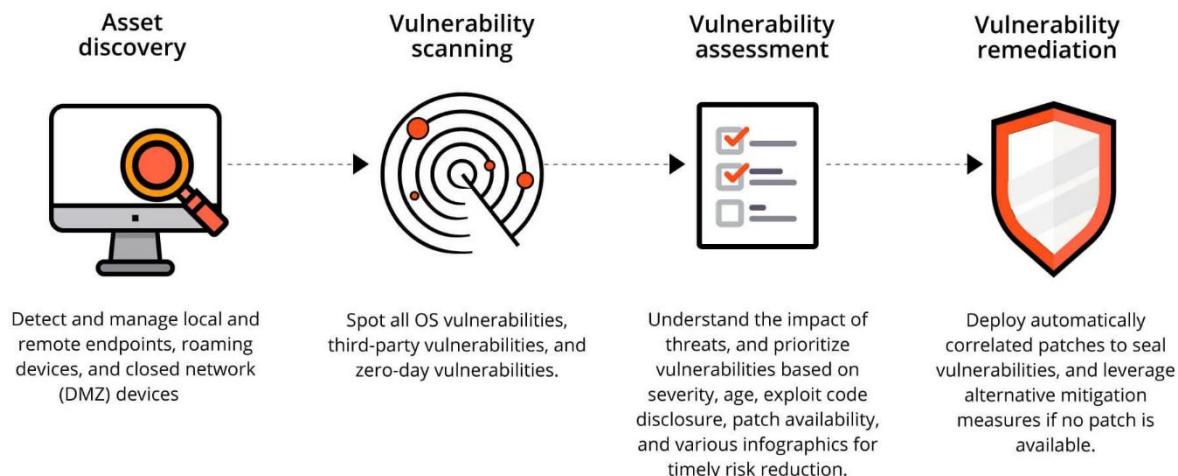
10. Secure Development Practices: Integrating security into the software development lifecycle helps ensure that applications deployed in the cloud are built with security in mind from the outset.



Key Practice



Vulnerability Assessment Tools for cloud



1. Nessus: Nessus is a widely used vulnerability scanner that can be deployed on-premises or in the cloud. It can assess cloud infrastructure for security vulnerabilities and compliance issues, providing detailed reports and recommendations for remediation.

2. Qualys Cloud Platform: Qualys offers a cloud-based platform for vulnerability management, including asset discovery, vulnerability scanning, and compliance monitoring. It supports various cloud environments and provides comprehensive security assessments.

3. Tenable.io: Tenable.io is a cloud-based vulnerability management platform that offers continuous visibility and assessment of cloud infrastructure, applications, and assets. It integrates with cloud providers and provides prioritized recommendations for risk reduction.

4. Rapid7 InsightVM: InsightVM is a cloud-based vulnerability assessment solution by Rapid7. It offers comprehensive visibility into cloud environments, with features such as asset discovery, vulnerability scanning, and risk prioritization.

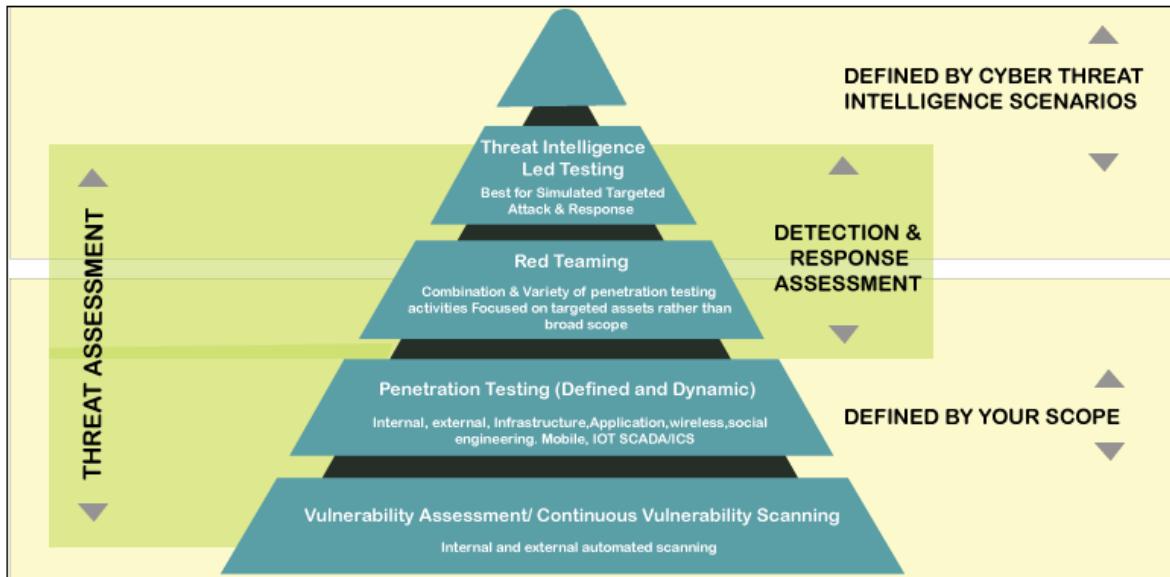
5. OpenVAS: OpenVAS (Open Vulnerability Assessment System) is an open-source vulnerability scanner that can be deployed in cloud environments. It provides scanning capabilities for identifying security vulnerabilities in cloud infrastructure and applications.

6. Amazon Inspector: Amazon Inspector is a security assessment service provided by AWS. It helps users improve the security and compliance of their AWS workloads by automatically assessing vulnerabilities and providing actionable recommendations.

7. Azure Security Center: Azure Security Center is a cloud security management service provided by Microsoft Azure. It offers vulnerability assessment capabilities for Azure resources, along with recommendations for improving security posture and compliance.

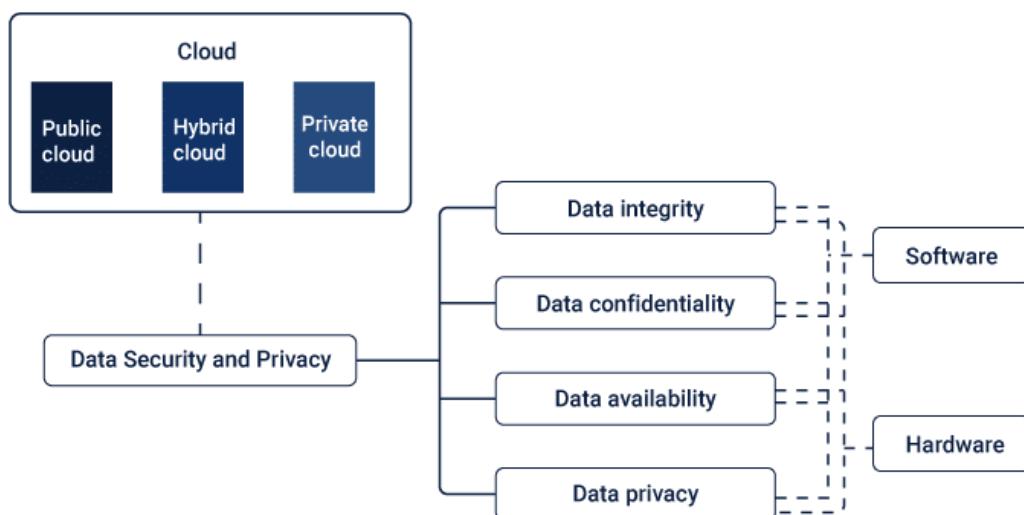
8. Google Cloud Security Command Center: Google Cloud Security Command Center provides security and risk visibility across Google Cloud Platform (GCP) services. It offers vulnerability scanning capabilities for identifying and remediating security risks in GCP resources.

Structure of Vulnerability Tools



Privacy and Security in cloud

Privacy and security are critical considerations when utilizing cloud computing services



TRUE ENGINEER

Data Integrity

In computing Data Integrity is vital. Any tampering, modification, or deletion of data can be very dangerous and costly. In a cloud computing services environment data can be easily lost or contaminated because of being accessed by unauthorized personnel or a systems breach. Better authentication and authorization protocols are needed to maintain the sanctity of the data. Two-factor authentication is one such protocol.

Data Availability

Data can be of various types. They can be structured or unstructured, in-transit or at-rest, or rarely used (as in back-up or disaster recovery data). Different service providers recognize this aspect of data and have tuned their cloud storage services to address this. The response time of data or its availability at the exact required moment is vital.

Cloud computing, which shares its storage environments among multiple clients, can sometimes create a situation where there is latency when retrieving the data of a particular client. But for critical processes, data may need to be stored in specified regions where they can be accessed quickly. This is a capability that the cloud computing services provider needs to build into its [cloud infrastructure](#).

Encryption

Today no data is sent across the internet or any open network without being encrypted. It is one of the mainstays of data privacy and offers the best protection against external threats. A plethora of Machine Learning algorithms and techniques are used to encrypt data to keep it from being compromised. But different types of data require different encryption and decryption techniques, and all such procedures need to be specifically requested from the service provider. It also needs to be spelled out clearly in the SLA.

Security Lapses by Authorised Personnel

Data can be compromised by employees, contractors, and other stakeholders due to carelessness or human error. Of course, it can also be done on purpose by malicious or disgruntled staff members. The staff members who have access to data, especially critical data, need to be trained and constantly made aware of the threat they pose to the organization's data. Perhaps, more security lapses happen because of employee carelessness than any other major external threat.

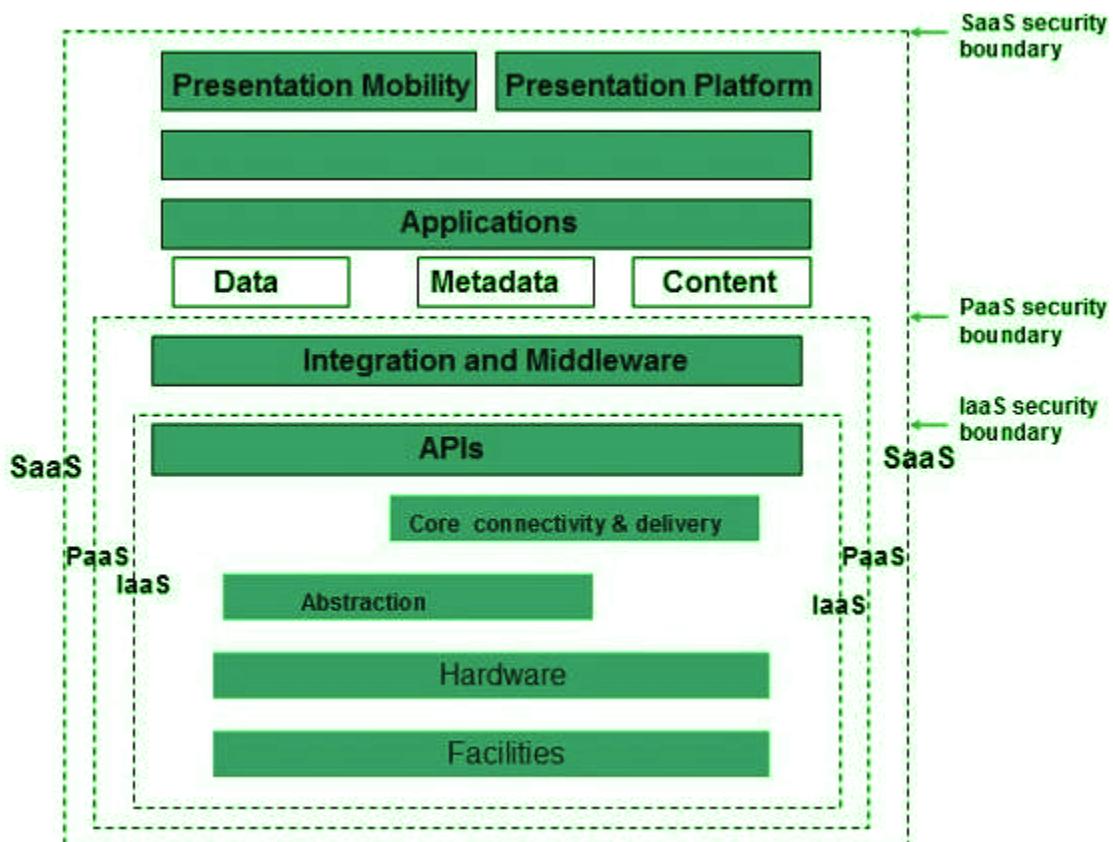
Threats

All threats are not necessarily data related. They can be because of infrastructural factors also: from compromised networks to sharing of host machines, from inelegance of the Virtual Machine (VM) to employee carelessness.

- Simple Breaches: There are many types of threats that originate from hacked accounts, lost passwords, open applications, or even virtual machine errors. Such breaches can be prevented through constant supervision and quick action is necessary to redress any problem. Regular backing-up of data, constant checks, and changing of access codes, etc. can prevent major calamities.
- Hijacked accounts, or even APIs, can cause much damage to the data in the cloud environment – it can be compromised very quickly without alerting any normal detection system.
- DoS Attacks: Denial of Service is a very common threat. If this happens, clients may not be able to access applications or data and can cause panic. Depending on the application or interface, a wrong message to a client showing Denial of Service can lead to embarrassment or even non-renewal of contracts in the long run. A system has to be up at all times for the client to feel confident about the service he/she is paying for.
- Malware Injections: These are scripts that run as valid SaaS programs inside the cloud. They can be used for contaminating data or disrupting performance, or worse.
- Network Security: The network is the most common doorway through which threats can enter. The data needs to be secure at all times, especially since it is traveling over the internet. Strong network traffic encryption techniques need to be put in place by the service provider to ensure less vulnerability on this front.

Cloud Computing Security Architecture

Cloud security architecture is often called cloud computing security architecture. It consists of security layers, design and structure of infrastructure, tools, software, platform, and best practices adopted within a cloud security solution. A cloud security architecture provides a visual and written model to establish how to secure and configure activities and operations in cloud; methods and controls in place for protection of applications, data; approach towards visibility in compliance, threats, and overall security posture.



Principles

- **Identification** – Overall cloud resource repository knowledge involving users, assets, business environment, policies, vulnerabilities, threats, risk management strategies which exist
- **Controls for security** – Parameters and policies implemented across users, assets, data, and infrastructure to manage overall security posture.
- **Security by design** – Standardized and repeated deployment of common use cases with security controls, standards, and audit requirements.
- **Compliance** – Integration of industry standard and regulatory standards into cloud architecture to meet the requirements.
- **Perimeter Security** – Management of connection points between corporate networks and public / external networks.
- **Segmentation** – To prevent lateral movement of attackers in cloud network segregation of sections.
- **User Identity and Access Management** – Visibility, understanding, and control on all users which have access to cloud assets. Access, permissions, and protocol enforcement.
- **Data Encryption** – Data at Rest and data in motion is encrypted to minimize breach impact.
- **Automation** – Rapid security and configuration provisioning and quick threat detection.
- **Logging and Monitoring** – activities are captured and monitored related to all connected systems and cloud-based services to ensure operations visibility, compliance, and early detection of threats.
- **Visibility in Multi-cloud** – Bring visibility in multiple cloud deployments by incorporating tools and processes.
- **Flexibility in Design** – Agility in architecture design to develop and incorporate new components and solutions without compromising security.

Cloud computing security Challenges

1. DDoS (Distributed Denial of Service):

- DDoS attacks aim to overwhelm a cloud service or infrastructure with a flood of traffic, rendering it inaccessible to legitimate users.
- Cloud providers often implement mitigation techniques such as traffic filtering, rate limiting, and distributed scrubbing centers to mitigate the impact of DDoS attacks. Additionally, organizations can deploy DDoS protection services and utilize Content Delivery Networks (CDNs) to distribute traffic geographically.

2. Data Breaches:

- Data breaches in the cloud can result from unauthorized access, misconfigurations, insider threats, or vulnerabilities in cloud services.
- Encryption of data at rest and in transit, strong access controls, monitoring and logging, regular security assessments, and compliance with data protection regulations (e.g., GDPR, HIPAA) help mitigate the risk of data breaches.

3. System Vulnerabilities:

- System vulnerabilities in cloud environments can arise from unpatched software, misconfigurations, insecure APIs, or shared resource vulnerabilities.
- Regular vulnerability assessments, patch management, secure configuration management, secure coding practices, and implementing security controls (e.g., firewalls, intrusion detection/prevention systems) help address system vulnerabilities.

4. Account Hijacking:

- Account hijacking involves unauthorized access to user accounts, often through techniques like phishing, credential theft, or weak authentication mechanisms.
- Implementing strong authentication mechanisms (e.g., multi-factor authentication), monitoring for suspicious account activities, user training on security best practices, and enforcing least privilege access controls help prevent account hijacking in the cloud.

TRUE ENGINEER

Misconfiguration

Cloud computing has emerged as a widely accepted approach for accessing resources remotely while simultaneously reducing costs. Cloud computing security concerns can be effectively mitigated through proper configuration of your cloud resources. Misconfiguration is the top cloud computing security challenge, as users must appropriately protect their data and applications in the cloud.

To prevent this cloud security threat, users must ensure their data is protected, and applications are configured correctly. It can be accomplished using a cloud storage service that offers security features such as encryption or access control. Additionally, implementing security measures such as authentication and password requirements can help protect sensitive data in the cloud. By taking these steps, users can increase the security of their cloud computing infrastructure and stay protected from cyber threats.

Unauthorized Access

Unauthorized access to data is one of the most common cloud security problems businesses face. The cloud provides a convenient way for companies to store and access data, which can make data vulnerable to cyber threats. Security and cloud computing threats can include unauthorized access to user data, theft of data, and malware attacks.

To protect their data from these threats, businesses must ensure that only authorized users can access it. Another security feature businesses can implement is encrypting sensitive data in the cloud. It will help ensure that only authorized users can access it. By implementing security measures such as encryption and backup procedures, businesses can safeguard their data from unauthorized access and ensure its integrity.

Hijacking of Accounts

Hijacking of user accounts is one of the major cloud security issues. Using cloud-based applications and services will increase the risk of account hijacking. As a result, users must be vigilant about protecting their passwords and other confidential information to stay secure in the cloud.

Users can protect themselves using strong passwords, security questions, and two-factor authentication to access their accounts. They can also monitor their account activity and take steps to protect themselves from unauthorized access or usage. This will help ensure that hackers cannot access their data or hijack their accounts. Overall, staying vigilant about security and updating your security measures are vital to the security of cloud computing.

Lack of Visibility

Cloud computing has made it easier for businesses to access and store their data online, but this convenience comes with risks. As a result, companies need to protect their data from unauthorized access and theft. However, cloud computing also poses security threats due to its reliance on remote servers. To ensure that their systems are vulnerable only to authorized sources, businesses must implement security measures such as strong authentication, data loss prevention (DLP), data breach detection, and data breach response.

With cloud computing, visibility is vital, and businesses must regularly audit security operations and procedures to detect vulnerabilities and threats before they become a real problem. By taking the necessary precautions and implementing security in cloud computing, organizations can ensure that their data remains secure in this cloud-based environment.

TRUE ENGINEER

Data Privacy/Confidentiality

Data privacy and confidentiality are critical issues when it comes to cloud computing. With cloud computing, businesses can access their data from anywhere worldwide, raising concerns about securing cloud computing. Companies don't have control over who can access their data, so they must ensure that only authorized users can access it. Data breaches can happen when hackers gain access to company data. In the coming years, there will be even more data privacy and confidentiality issues due to the rise of big data and the increased use of cloud computing in business.

Data privacy and confidentiality issues will continue to be essential concerns for businesses in the years ahead as data-intensive applications grow in popularity. [Managed IT Services Charlotte](#) experts help to ensure proper security measures and data practice for a cloud-ready organization to avoid data breach risks.

External Sharing of Data

External data sharing is one of the leading security issues in cloud computing that businesses face. This issue arises when data is shared with third-party providers who must be vetted and approved by the organization. As a result, external data sharing can lead to the loss of critical business information and theft and fraud. To prevent these issues in cloud security, companies must implement robust security measures, such as encryption and data management practices. In addition, it will help ensure that sensitive data remains secure and confidential.

By implementing appropriate security measures, companies can protect their data from unauthorized access and ensure its reliability and integrity. Overall, external data sharing is a major cloud security concern that businesses must address to stay ahead of the competition.

Unsecure Third-party Resources

Third-party resources are applications, websites, and services outside the cloud provider's control. These resources may have cloud security vulnerabilities, and unauthorized access to your data is possible. Additionally, unsecured third-party resources may allow hackers to access your cloud data. These vulnerabilities can put your security at risk. Therefore, ensuring that only trusted, secure resources are used for cloud computing is essential. In addition, it will help ensure that only authorized individuals access data and reduce the risk of unauthorized data loss or breach.

Unsecured third-party resources can pose a threat to cloud security, especially when interacting with sensitive data in cloud storage accounts. Hackers can access these resources to gain access to your cloud data and systems. Implementing strong security controls such as multi-factor authentication and enforcing strict password policies can help safeguard against this risk. In addition, by restricting access to only trusted resources, you can ensure that only authorized individuals access data and reduce the risk of unauthorized data loss or breach.

Virtualization security management

Virtualized security, also known as security virtualization, offers dynamic and flexible security solutions tailored for virtualized IT environments.

1. Software-based Security Solutions:

- Virtualized security solutions are software-based and designed to operate within virtualized IT environments. They can be deployed as virtual appliances, software-defined security services, or cloud-based security platforms.
- Unlike traditional hardware-based security appliances, virtualized security solutions are not tied to specific physical devices, offering greater flexibility and scalability.

2. Dynamic Deployment and Mobility:

- Virtualized security solutions can be dynamically deployed and moved within the network to adapt to changing demands and workloads. This flexibility is crucial in virtualized environments where workloads are spun up, moved, and scaled dynamically.
- Security services can follow and protect workloads as they move across virtualized infrastructure, ensuring consistent security regardless of workload location or changes in the network topology.

3. Cloud Compatibility:

- Virtualized security solutions are well-suited for cloud environments, including public, private, and hybrid clouds. They can be seamlessly integrated with cloud platforms and services, providing security controls that align with cloud security considerations.
- Virtualized security enables isolation and segmentation of multitenant environments in public clouds, ensuring that each tenant's resources are adequately protected and isolated from others.

4. Flexibility for Hybrid and Multi-Cloud Environments:

- In hybrid and multi-cloud environments, where workloads and data span across multiple cloud platforms and vendors, virtualized security offers the flexibility to provide consistent security controls and policies.
- Security functions can be deployed uniformly across different cloud environments, ensuring a cohesive security posture and simplifying management across the entire cloud ecosystem.

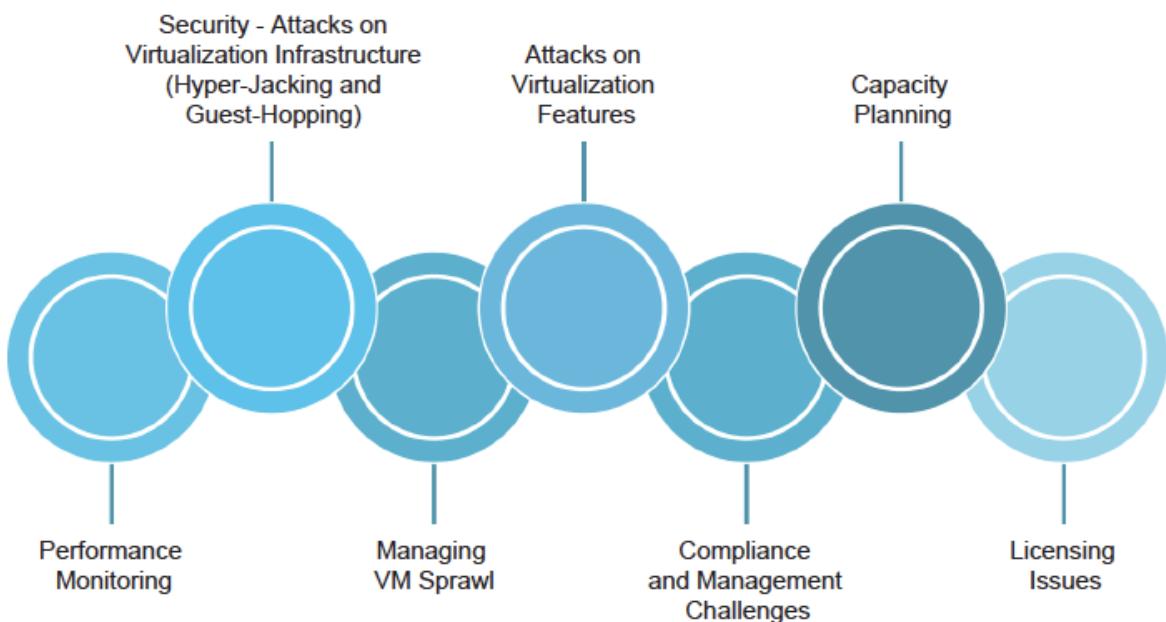
5. Scalability and Cost-effectiveness:

- Virtualized security solutions can scale up or down based on demand, allowing organizations to adjust their security posture as needed without being constrained by hardware limitations.
- Virtualized security often offers cost savings compared to traditional hardware-based security, as it eliminates the need for expensive physical appliances and allows for efficient resource utilization through virtualization technology.

What are the benefits of virtualized security?

Virtualized security is now effectively necessary to keep up with the complex security demands of a virtualized network, plus it's more flexible and efficient than traditional physical security. Here are some of its specific benefits:

- **Cost-effectiveness:** Virtualized security allows an enterprise to maintain a secure network without a large increase in spending on expensive proprietary hardware. Pricing for cloud-based virtualized security services is often determined by usage, which can mean additional savings for organizations that use resources efficiently.
- **Flexibility:** Virtualized security functions can follow workloads anywhere, which is crucial in a virtualized environment. It provides protection across multiple data centers and in multi-cloud and hybrid cloud environments, allowing an organization to take advantage of the full benefits of virtualization while also keeping data secure.
- **Operational efficiency:** Quicker and easier to deploy than hardware-based security, virtualized security doesn't require IT teams to set up and configure multiple hardware appliances. Instead, they can set up security systems through centralized software, enabling rapid scaling. Using software to run security technology also allows security tasks to be automated, freeing up additional time for IT teams.
- **Regulatory compliance:** Traditional hardware-based security is static and unable to keep up with the demands of a virtualized network, making virtualized security a necessity for organizations that need to maintain regulatory compliance.



Secure Execution Environments

Secure execution environments provide isolated and protected spaces within computing systems where sensitive operations, such as running critical applications or processing confidential data, can be performed securely. These environments are designed to protect against various threats, including unauthorized access, malware, data breaches, and tampering. Here are some key aspects of secure execution environments:

1. Isolation:

- Secure execution environments utilize isolation mechanisms to separate and protect sensitive operations from other parts of the system. This isolation prevents unauthorized access and interference by malicious actors or software.

2. Hardware-based Security Features:

- Many secure execution environments leverage hardware-based security features provided by modern processors and chipsets. These features may include hardware-enforced memory protection, secure boot, and trusted execution environments (TEEs) such as Intel SGX or ARM TrustZone.

3. Encryption and Secure Storage:

- Secure execution environments often incorporate encryption and secure storage mechanisms to protect sensitive data both at rest and in transit. Encryption ensures that data remains confidential even if it is accessed by unauthorized parties.

4. Access Controls:

- Access controls are implemented to restrict and manage the permissions granted within the secure execution environment. This includes authentication mechanisms, role-based access control (RBAC), and least privilege principles to limit access to authorized users and processes.

5. Integrity Verification:

- Secure execution environments may employ integrity verification mechanisms to ensure the integrity of code, data, and configurations within the environment. This helps detect and prevent unauthorized modifications or tampering.

6. Secure Communication Channels:

- Secure execution environments establish encrypted communication channels for exchanging data and instructions between trusted components. This protects against eavesdropping, interception, and tampering of communication channels.

7. Secure Development Practices:

- Secure execution environments are built using secure development practices, including secure coding, vulnerability assessments, and security testing. This helps minimize the risk of security vulnerabilities and weaknesses in the environment.

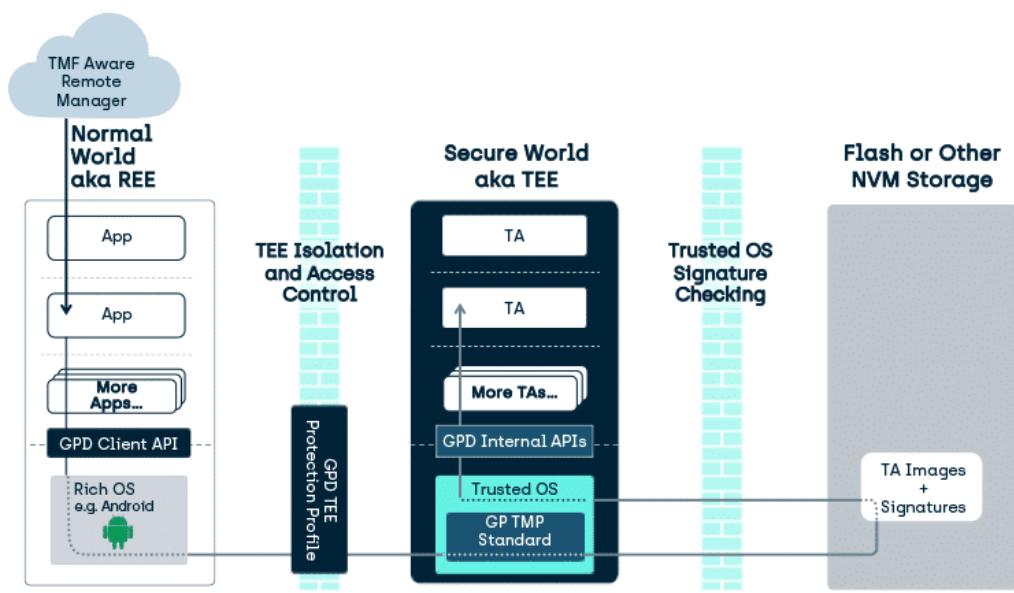
TRUE ENGINEER

8. Auditing and Logging:

- Auditing and logging mechanisms are implemented to record and monitor activities within the secure execution environment. This enables accountability, forensic analysis, and detection of security incidents or policy violations.

9. Continuous Monitoring and Threat Detection:

- Secure execution environments are continuously monitored for signs of security threats or anomalies. Intrusion detection systems (IDS), anomaly detection algorithms, and behavioral analysis help detect and respond to security incidents in real-time.



Unit 05

QOS Issues in Cloud (Quality Of Service)

Quality of Service (QoS) issues in cloud computing can arise due to various factors that affect the performance, reliability, and availability of cloud services. Here are some common QoS issues in cloud computing:

1. Network Latency:

- Network latency refers to the delay in data transmission between client devices and cloud servers. High network latency can degrade the responsiveness of cloud applications and impact user experience, particularly for real-time applications like video streaming or online gaming.

2. Bandwidth Limitations:

- Bandwidth limitations can restrict the amount of data that can be transmitted between clients and cloud servers. Insufficient bandwidth can lead to slow data transfers, buffering, and degraded performance for data-intensive applications.

3. Resource Contention:

- Resource contention occurs when multiple users or applications compete for resources within a shared cloud infrastructure. This can lead to performance degradation, unpredictable response times, and bottlenecks, especially during peak usage periods.

4. Service Outages:

- Service outages, whether due to hardware failures, network disruptions, or software issues, can result in downtime and unavailability of cloud services. Lack of redundancy, inadequate failover mechanisms, and insufficient disaster recovery planning can exacerbate the impact of service outages.

TRUE ENGINEER

5. Inadequate Scalability:

- Inadequate scalability can limit the ability of cloud services to handle fluctuations in workload demands. If cloud resources cannot scale dynamically to accommodate increased traffic or computational needs, performance may suffer, and users may experience slowdowns or service interruptions.

6. Security Concerns:

- Security issues such as data breaches, unauthorized access, or malware attacks can compromise the confidentiality, integrity, and availability of cloud services. Implementing robust security controls, encryption, access controls, and security monitoring helps mitigate security-related QoS issues.

7. SLA Violations:

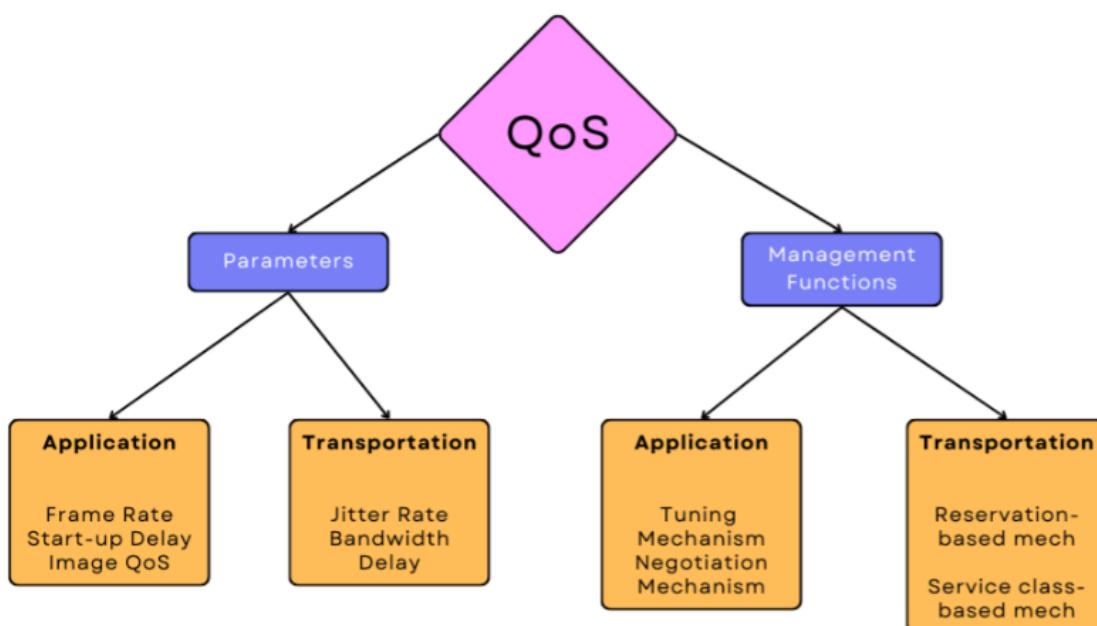
- Service Level Agreement (SLA) violations occur when cloud service providers fail to meet the performance, uptime, or availability guarantees specified in their SLAs. SLA breaches can result in financial penalties, reputational damage, and loss of customer trust.

8. Geographical Latency:

- Geographical latency refers to the delay introduced by the physical distance between clients and cloud data centers. Users located far from cloud regions may experience higher latency, affecting the responsiveness of cloud applications and services.

9. Vendor Lock-In:

- Vendor lock-in occurs when organizations become dependent on a single cloud service provider, limiting their ability to switch providers or migrate workloads easily. Lack of interoperability and portability can hinder flexibility and agility in managing cloud services.



Data Migration



Data migration refers to the process of transferring data from one storage system or environment to another. This can involve moving data between different types of storage infrastructure, such as from on-premises servers to cloud storage, between cloud providers, or within the same cloud environment.

Organizations can initiate a data migration process due to various reasons:

- To replace or upgrade servers or storage infrastructure
- To move on-premises infrastructure to cloud-based platforms
- To move data between third-party cloud services (cloud data migration)
- To perform infrastructure maintenance
- To consolidate websites
- To migrate databases or applications
- To move data during a data center relocation or a merger
- To install software upgrades

Why is data migration important?

Data is king for modern business competitiveness. Companies rely on proper data management to procure services, manage business processes, and ensure customer satisfaction, business continuity, and a steady revenue stream.

Data migration ensures that critical organizational data is securely transferred to another app, storage system, or the cloud. Migrating data between platforms can be challenging but provides companies with many benefits. A successful migration process can boost productivity, reduce storage costs, upgrade applications and services, and more.

What are the different types of data migration?

There are six primary types of data migration. Let's explore them below.

Storage migration

Storage migration refers to moving data from one computer storage system to another. The process often involves physical data migration from one hardware storage system to a destination system.

Storage migration is commonly done to upgrade existing storage equipment to a more sophisticated storage infrastructure. In such a scenario, the process involves moving data from an old to a new system - paper to digital, tape drives to HDDs, HDD to solid-state drives (SSD), and physical storage to a cloud computing environment (virtual storage).

Often, storage migration is not driven by insufficient storage space but rather by the need to upgrade storage technology. Typically, this approach doesn't alter or format the data. However, companies can use the opportunity to perform data validation and reduction by detecting corrupt or obsolete data.

Application data migration

Application migration (or "app migration") involves transferring data from an app or program from one computing environment to another. The process usually occurs when a company changes application software or switches to another application vendor. If the new application requires different application interactions, the application migration may require radical data transformations.

A significant challenge for application migration comes from the source and target systems having specific data models and using different data formats. Vendors can provide application programming interfaces (APIs) to protect data integrity. Moreover, organizations may benefit from vendor web interfaces to facilitate application migration and middleware to fill the gaps between the app and operating systems.

Database migration

Databases house and structure data in an organized way to enable more efficient storage technologies. Databases are managed via database management systems (DBMS).

Database migration involves moving data from one database management system to another or upgrading from an old DBMS version to the latest one for the same DBMS. The former scenario is more challenging as the source and target systems often use different data structures.

Database migration is commonly done when a company changes database vendors, moves the database to the cloud, or upgrades the database software. It's critical to back up all databases before migration.

Business process migration

Business process migration involves moving business application data and data regarding business processes and metrics to a new environment. The business process metrics commonly include product, customer, and operational data.

A common reason for business process migration is business optimization and reorganization or mergers and acquisitions (M&A). This approach to data migration is necessary for many organizations to enter new markets and stay competitive in an ever-evolving field.

Cloud migration

Cloud migration has become a common form of data migration. **Cloud migration** refers to moving data or applications from local (on-premises) storage to the cloud or from one cloud platform to another. (cloud storage migration)

The cloud environment provides on-demand scalability and flexibility and reduces capital expenditure (CapEx) for on-premises infrastructures. Cloud service providers offer various features regarding your storage, application, database, and cloud migration needs.

Data center migration

A data center houses an organization's data storage infrastructure required to maintain critical applications. The data center comprises network routers, servers, computers, switches, storage devices, and all related data equipment.

Data center migration refers to migrating data from an on-premises data center to a new physical location or a new system (the cloud) or from old data center infrastructure to new infrastructure equipment at the same (physical) location.

Streaming in Cloud

Streaming in the cloud refers to the delivery of multimedia content, such as audio, video, or live broadcasts, over the internet from cloud-based servers to end-users' devices in real-time or near real-time. Cloud-based streaming services offer scalability, flexibility, and cost-effectiveness compared to traditional on-premises infrastructure.



TRUE ENGINEER

1. Content Storage and Encoding:

- Multimedia content is typically stored in the cloud in formats optimized for streaming delivery. This may involve encoding the content into different bitrates and resolutions to accommodate varying network conditions and device capabilities.
- Cloud storage services, such as Amazon S3, Google Cloud Storage, or Azure Blob Storage, are commonly used to store multimedia files securely and durably.

2. Content Delivery Networks (CDNs):

- CDNs are networks of distributed servers located in multiple geographic locations. They cache and deliver multimedia content to end-users from servers located closer to their geographical location, reducing latency and improving streaming performance.
- Cloud-based CDNs, such as Amazon CloudFront, Google Cloud CDN, or Azure CDN, integrate seamlessly with cloud streaming services to deliver high-quality streaming experiences to users worldwide.

3. Streaming Protocols and Formats:

- Various streaming protocols and formats are used to deliver multimedia content over the internet. Common protocols include HTTP Live Streaming (HLS), Dynamic Adaptive Streaming over HTTP (DASH), and Real-Time Messaging Protocol (RTMP).
- Cloud streaming platforms often support multiple streaming protocols and formats to ensure compatibility with a wide range of devices and players.

4. Scalability and Elasticity:

- Cloud-based streaming services offer scalability and elasticity to handle fluctuations in streaming demand. They can automatically scale up or down based on traffic spikes or changes in viewership, ensuring a seamless streaming experience for users without the need for additional infrastructure provisioning.

5. Security and Digital Rights Management (DRM):

- Security is a critical consideration for cloud-based streaming services to protect content from unauthorized access, piracy, and content theft.
- DRM solutions, such as encryption, access controls, and watermarking, are implemented to secure streaming content and enforce digital rights management policies.

6. Analytics and Monitoring:

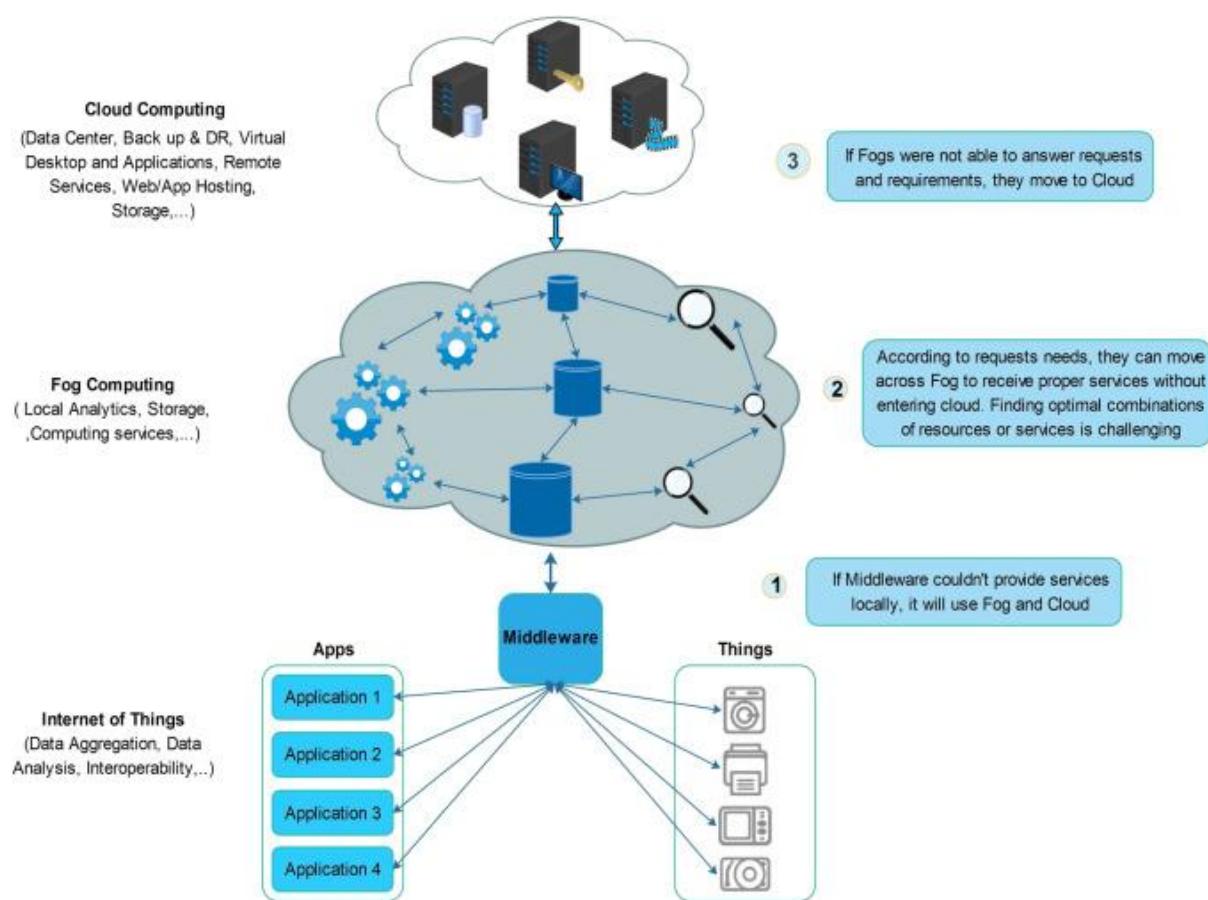
- Cloud streaming platforms provide analytics and monitoring capabilities to track streaming performance, audience engagement, and viewer behavior.
- Real-time analytics dashboards and reports enable content providers to gain insights into streaming metrics, audience demographics, and content consumption patterns, helping them optimize content delivery and monetization strategies.

7. Integration with Streaming Platforms:

- Cloud-based streaming services integrate with streaming platforms, content management systems (CMS), and monetization platforms to streamline content publishing, distribution, and monetization workflows.
- APIs and SDKs provided by cloud streaming platforms facilitate seamless integration with third-party services and applications, enabling customized streaming solutions and enhanced user experiences.

Cloud Middleware

Cloud middleware refers to software components or services deployed in cloud environments to facilitate communication, integration, and interaction between different applications, services, or components. Middleware acts as an intermediary layer between software components, abstracting the complexities of underlying infrastructure and providing standardized interfaces and protocols for interoperability.



1. Integration and Interoperability:

- Cloud middleware enables seamless integration and interoperability between heterogeneous systems, applications, and services deployed across distributed cloud environments.
- It provides standardized communication protocols, message formats, and data exchange mechanisms to facilitate interaction between diverse software components.

TRUE ENGINEER

2. Service Orchestration:

- Middleware platforms in the cloud support service orchestration and choreography, enabling the composition and coordination of distributed services to automate business processes and workflows.
- They provide tools and frameworks for defining, managing, and executing service workflows, including service discovery, invocation, and monitoring.

3. Message Brokering and Event Processing:

- Cloud middleware includes messaging systems and event processing engines that facilitate asynchronous communication, pub/sub messaging, and event-driven architectures.
- Message brokers, such as Apache Kafka, RabbitMQ, or Amazon SQS, enable reliable and scalable message delivery between distributed components in the cloud.

4. Data Integration and Transformation:

- Middleware platforms in the cloud support data integration and transformation capabilities to enable seamless data exchange and synchronization between disparate systems and data sources.
- They provide tools for data mapping, transformation, cleansing, and replication, as well as support for various data formats and protocols.

5. API Management:

- Cloud middleware includes API management platforms that provide capabilities for creating, publishing, securing, and managing APIs to expose backend services and functionalities to external consumers.
- API gateways, API portals, rate limiting, authentication, and monitoring features are commonly included in API management solutions.

6. Security and Compliance:

- Middleware platforms in the cloud include security features and controls to ensure data confidentiality, integrity, and availability.
- They support identity and access management, encryption, authentication, authorization, audit logging, and compliance monitoring to protect sensitive data and comply with regulatory requirements.

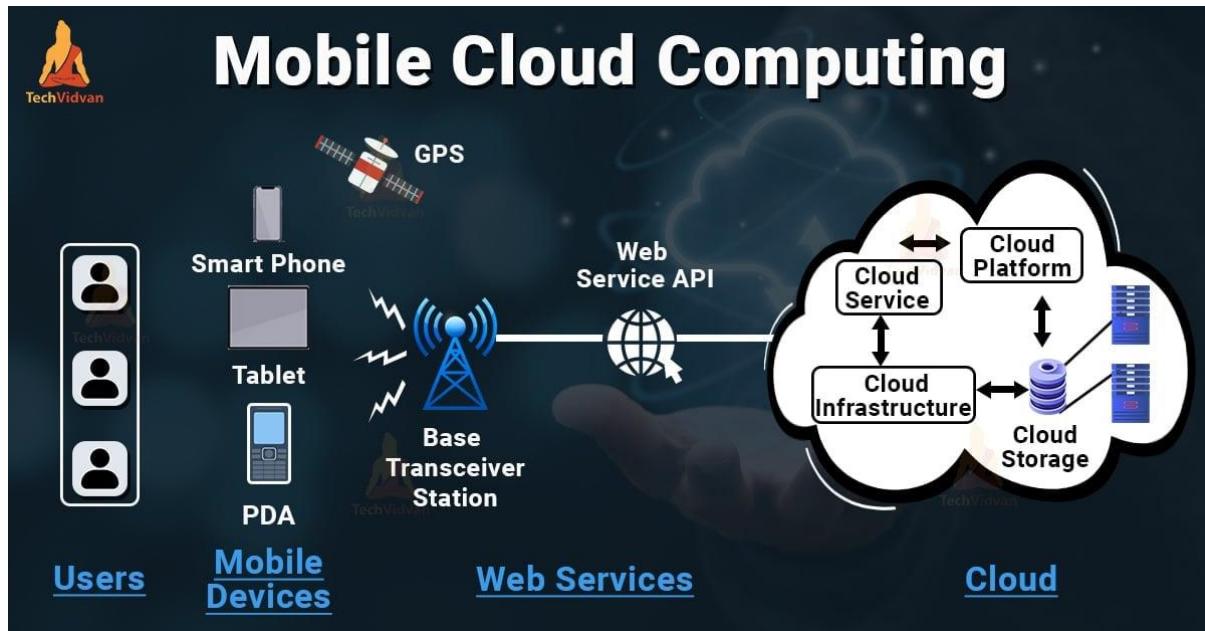
7. Scalability and Performance:

- Cloud middleware is designed to be scalable and performant, capable of handling large volumes of transactions, requests, and data streams across distributed cloud environments.
- It leverages cloud-native architectures, microservices, containerization, and auto-scaling capabilities to dynamically scale resources based on demand.

8. Monitoring and Management:

- Cloud middleware platforms provide monitoring, management, and diagnostics tools to monitor the health, performance, and availability of distributed applications and services.
- They offer dashboards, alerts, logging, and analytics capabilities to track and analyze system metrics and troubleshoot issues in real-time.

Mobile Cloud Computing



Mobile cloud computing (MCC) combines the capabilities of mobile devices and cloud computing to provide enhanced computational resources, storage, and services to mobile users. It leverages the scalability, flexibility, and accessibility of cloud computing to overcome the limitations of mobile devices, such as limited processing power, storage, and battery life.

1. Offloading Computational Tasks:

- Mobile cloud computing enables offloading computationally intensive tasks from mobile devices to remote cloud servers. Tasks such as complex calculations, data processing, image/video rendering, and machine learning inference can be performed in the cloud, reducing the burden on mobile devices and improving performance.

2. Data Storage and Synchronization:

- Cloud storage services provide mobile users with scalable and accessible storage for their data, files, and media content. Mobile apps can store data in the cloud, enabling seamless synchronization and access from multiple devices.
- Cloud storage also facilitates data backup, sharing, and collaboration among users, enhancing data management and accessibility on mobile devices.

TRUE ENGINEER

3. Augmented Computing Power:

- Mobile cloud computing augments the computing power of mobile devices by leveraging cloud resources for complex computations and analytics. This enables advanced functionalities and applications on mobile devices, such as augmented reality (AR), virtual reality (VR), and real-time gaming.

4. Scalability and Flexibility:

- Cloud computing offers scalability and flexibility to mobile applications and services, allowing them to handle variable workloads, user demands, and resource requirements.
- Mobile apps can dynamically scale their computing and storage resources based on demand, leveraging cloud infrastructure and services to meet changing user needs.

5. Location Independence:

- Mobile cloud computing enables location-independent access to computing resources and services. Mobile users can access cloud-based applications, data, and services from anywhere with internet connectivity, regardless of their physical location.
- This flexibility allows users to work, collaborate, and interact with digital content on their mobile devices while on the go.

6. Cost-effectiveness:

- Mobile cloud computing can be cost-effective for both users and developers. Cloud services often offer pay-as-you-go pricing models, allowing users to pay only for the resources and services they consume.
- Developers can leverage cloud infrastructure and platforms to build and deploy mobile applications more efficiently, reducing development costs and time-to-market.

7. Security and Privacy:

- Security and privacy are important considerations in mobile cloud computing. Cloud providers implement security controls, encryption, access controls, and compliance measures to protect data and ensure user privacy.
- Mobile apps and devices must also implement security best practices, such as secure communication protocols, authentication, and data encryption, to protect sensitive information and prevent unauthorized access.

Features of Mobile Cloud Computing

- It helps in the rapid development and shared resources of Mobile Applications.
- MCC also supports multiple development techniques and various devices.
- It helps to improve reliability and helps to keep the information secure by keeping a backup and storing the data on the Cloud.
- Applications utilize the device's resources at a lower rate because they have the support of Cloud Technology.
- Mobile devices are connected to the services and are delivered on the Application Programming Interface (API) architecture.

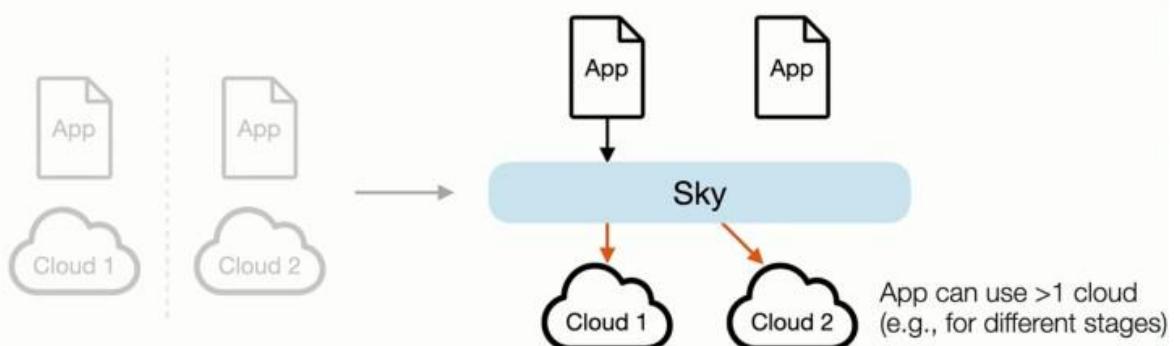
Sky computing

Cloud computing has transformed the way businesses and individuals leverage computing power but its true potential has been limited by the fragmentation caused by diverse services and APIs offered by different cloud providers. However, a paradigm shift is underway with the emergence of Sky Computing that addresses the challenges of fragmentation and aims to achieve true utility computing. Sky Computing envisions a compatibility layer that enables applications developed for one cloud to seamlessly run on different clouds without modifications. It also introduces an inter-cloud layer that abstracts away the underlying clouds, enabling applications to run without specific cloud awareness.

Sky computing emerges as a solution to address the challenges of fragmentation and achieve the ultimate goal of true utility computing. It envisions a compatibility layer that enables applications developed for one cloud to run seamlessly on different clouds without requiring modifications.

Sky Computing

Enabling transparent multicloud [Stoica & Shenker, HotOS '21]



Sky Computing Benefits

- ① Access to new capabilities from different clouds
 - ② Enhanced security through distributed trust
 - ③ Improved reliability of systems
 - ④ Better performance through resource aggregation
- 💰 Potential cost savings with optimal cloud selection

